

# **Stanford CS224W:** **Graph Neural Networks**

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>



# ANNOUNCEMENTS

- **Today (10/07):** Colab 1 due, Colab 2 out
- **Next Thursday (10/14):** HW 1 due, HW 2 out
- **Project proposals due on Tuesday 10/19**
  - If you are looking for project partners, check out / add yourself to our pinned Ed post ("Project Partner Thread")  
-- reach out to each other!
  - We strongly encourage groups of 3, but groups of 1 or 2 are allowed

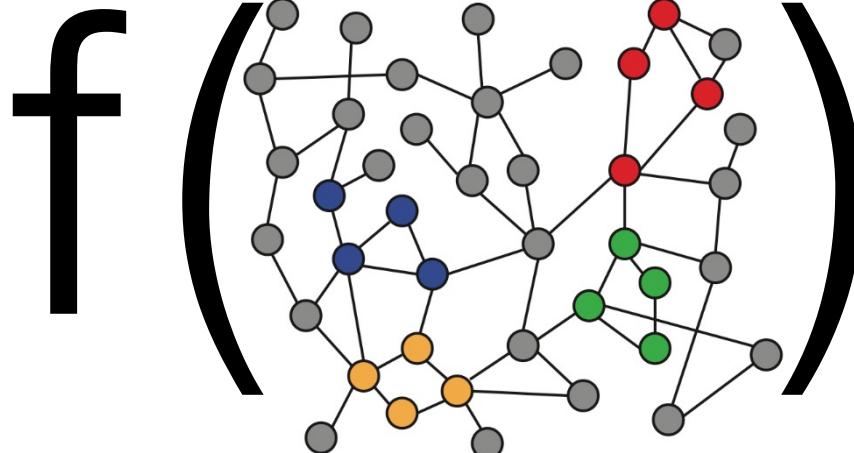
CS224W: Machine Learning with Graphs  
Jure Leskovec, Stanford University  
<http://cs224w.stanford.edu>



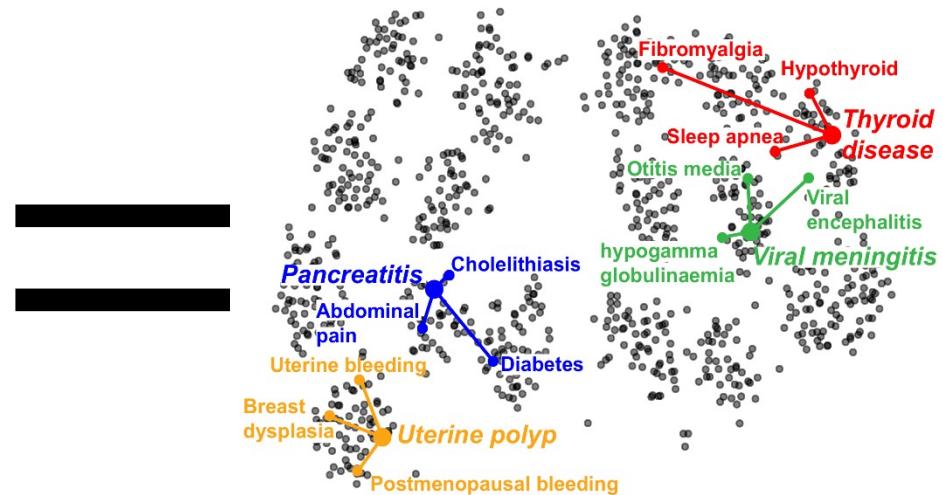
# Recap: Node Embeddings

图嵌入表示学习

- **Intuition:** Map nodes to  $d$ -dimensional embeddings such that similar nodes in the graph are embedded close together



Input graph



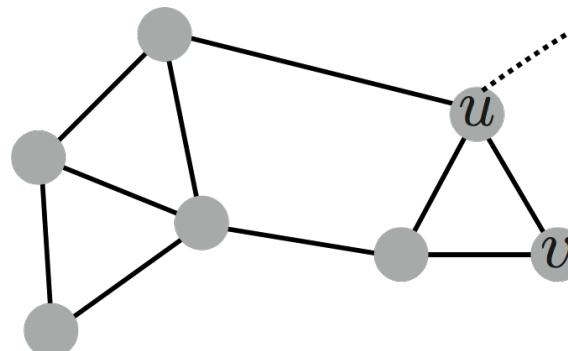
2D node embeddings

How to learn mapping function  $f$ ?

# Recap: Node Embeddings

Goal:  $\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$

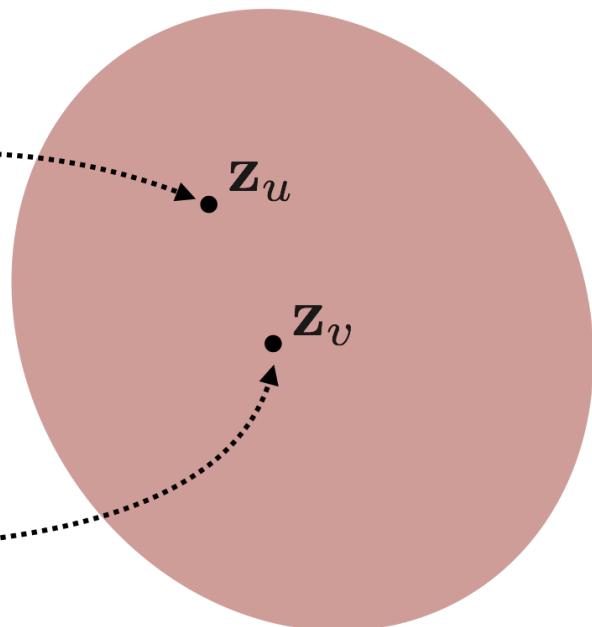
Need to define!



encode nodes

$\text{ENC}(u)$

$\text{ENC}(v)$



Input network

$d$ -dimensional  
embedding space

# Recap: Two Key Components

- **Encoder:** Maps each node to a low-dimensional vector

$\text{ENC}(v) = \mathbf{z}_v$

*d*-dimensional embedding

node in the input graph

- **Similarity function:** Specifies how the relationships in vector space map to the relationships in the original network

$$\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$$

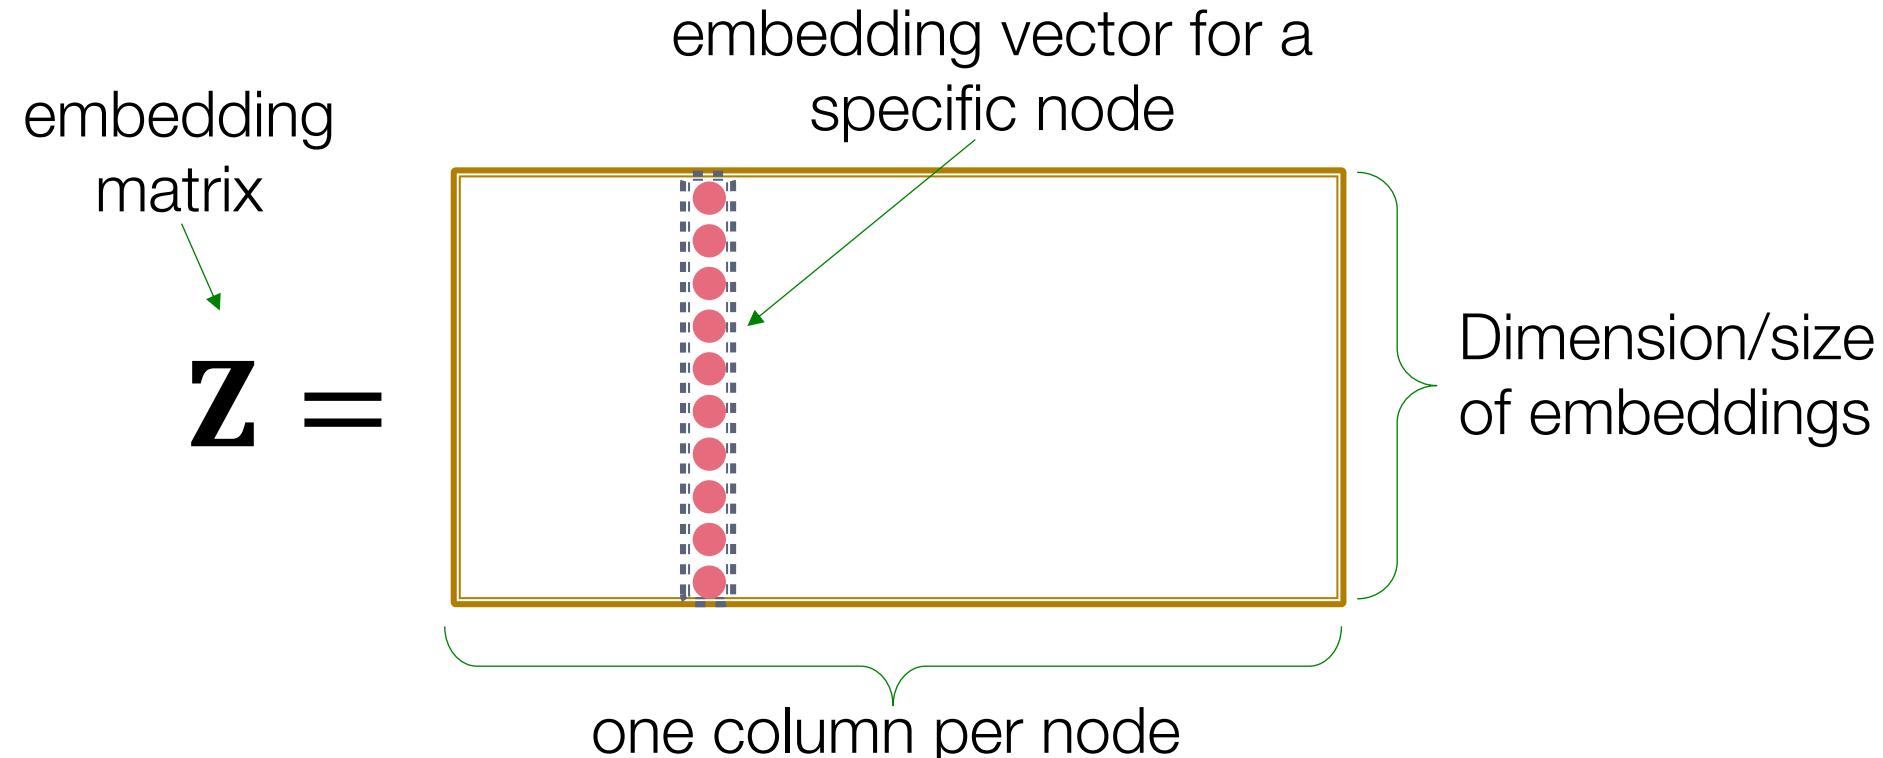
Similarity of  $u$  and  $v$  in  
the original network

**Decoder**

dot product between node embeddings

# Recap: “Shallow” Encoding

Simplest encoding approach: **Encoder is just an embedding-lookup**



# Recap: Shallow Encoders

DeepWalk, Node2vec, LINE

- Limitations of shallow embedding methods:
  - **$O(|V|)$  parameters are needed:** 每个节点的嵌入向量都需单独训练  
参数共享
    - No sharing of parameters between nodes
    - Every node has its own unique embedding
  - **Inherently “transductive”:** 直推式 无法泛化到新图 / 新节点
    - Cannot generate embeddings for nodes that are not seen during training
  - **Do not incorporate node features:** 没有用到节点属性特征标注
    - Nodes in many graphs have features that we can and should leverage  
利用

# Today: Deep Graph Encoders

图深度学习 图神经网络

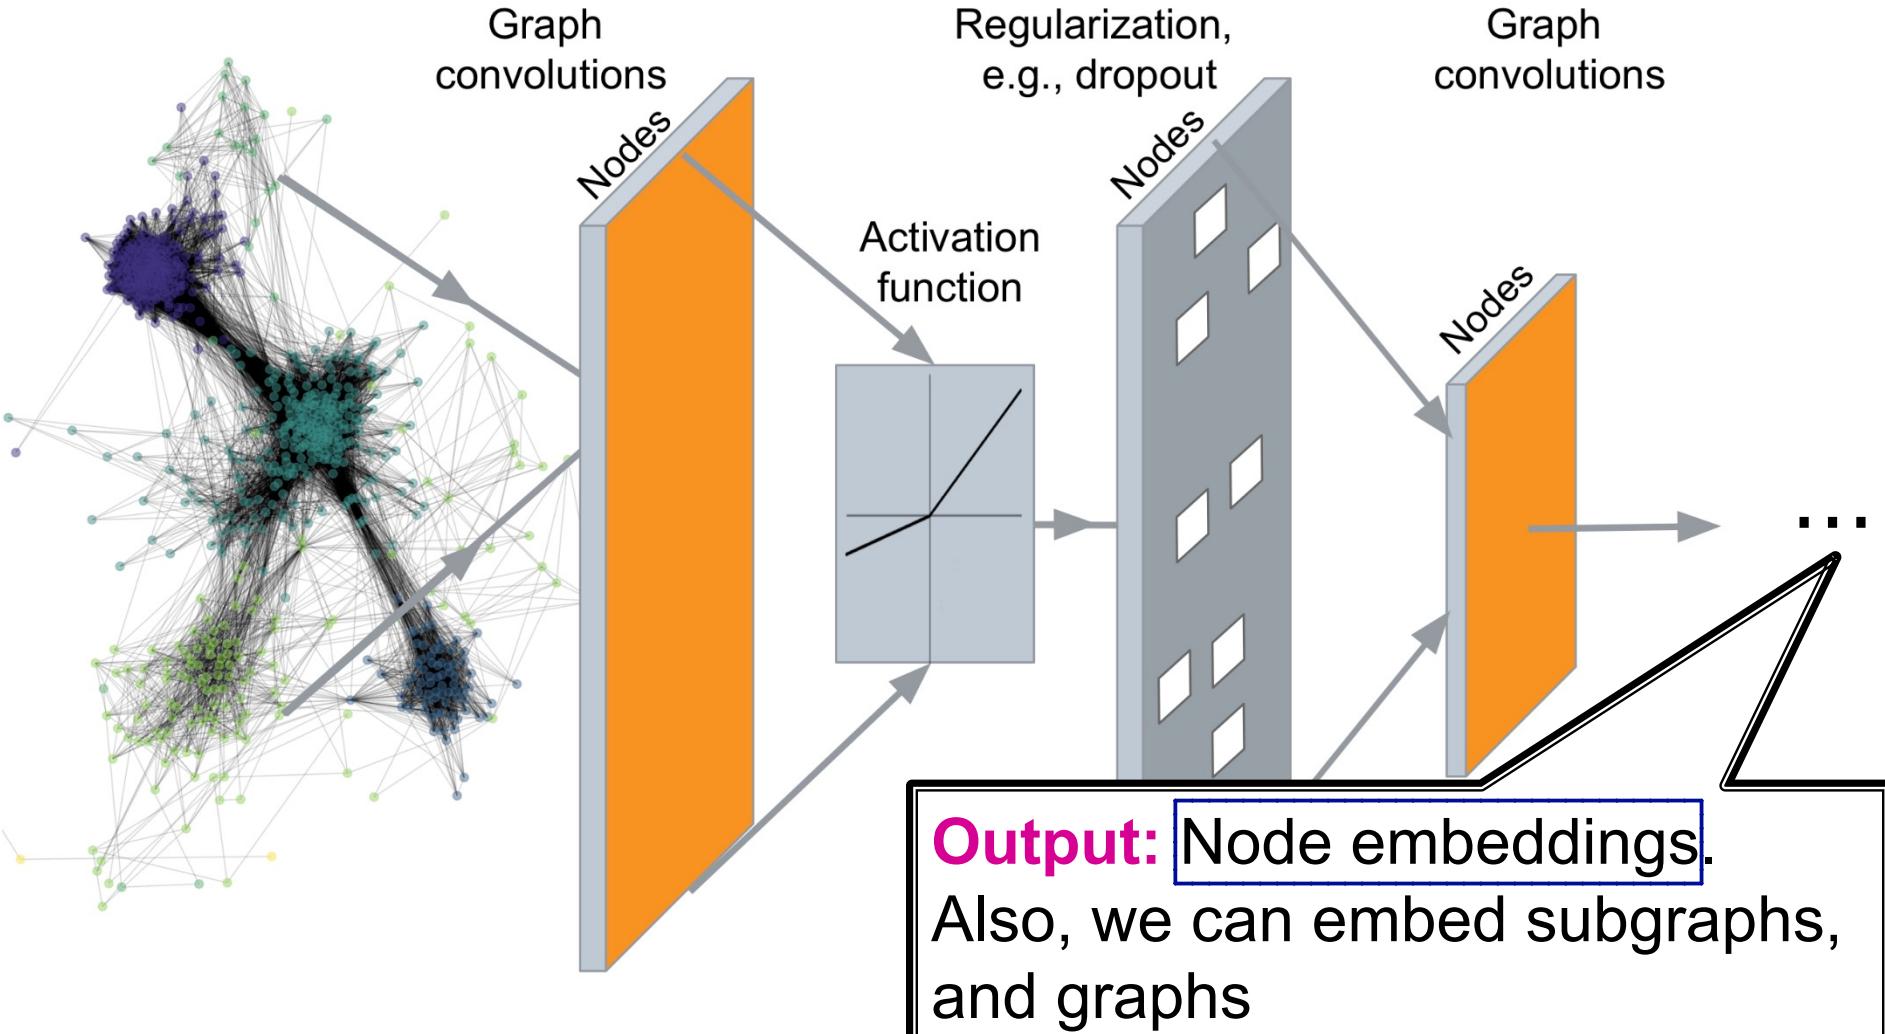
- **Today:** We will now discuss deep learning methods based on **graph neural networks (GNNs)**:

$$\text{ENC}(\nu) =$$

multiple layers of  
non-linear transformations  
based on graph structure

- **Note:** All these deep encoders can be **combined with node similarity functions** defined in the Lecture 3.

# Deep Graph Encoders

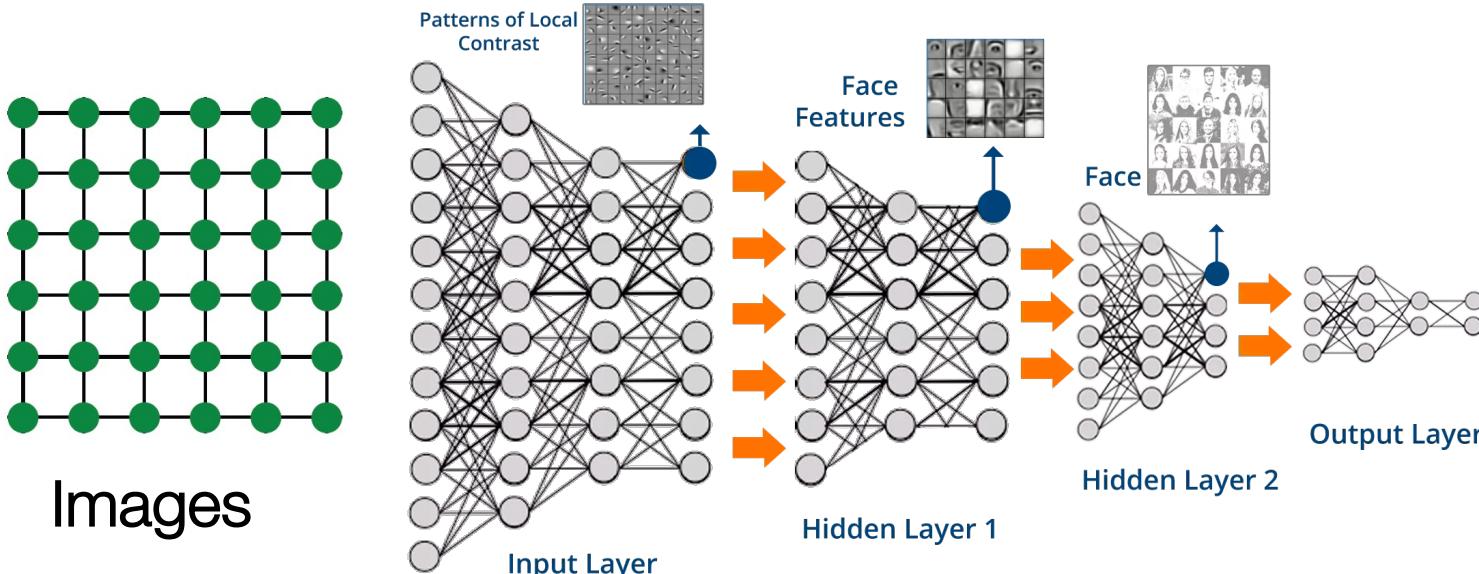


# Tasks on Networks

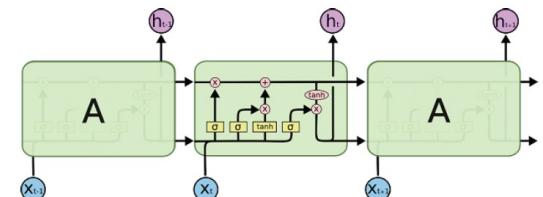
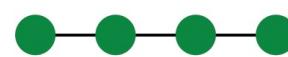
## Tasks we will be able to solve:

- Node classification
  - Predict a type of a given node
- Link prediction
  - Predict whether two nodes are linked
- Community detection
  - Identify densely linked clusters of nodes
- Network similarity
  - How similar are two (sub)networks

# Modern ML Toolbox



Text/Speech

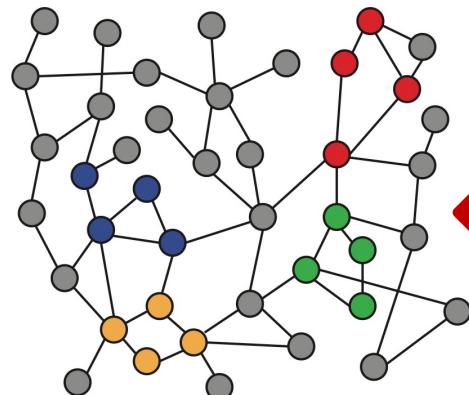


Modern deep learning toolbox is designed  
for simple sequences & grids

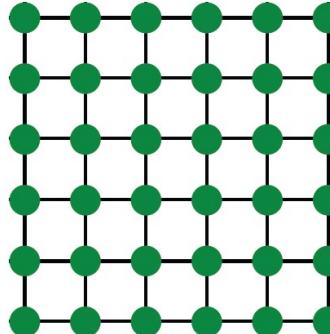
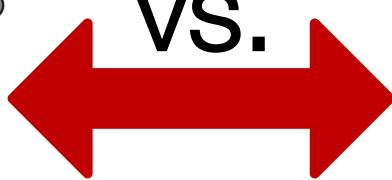
# Why is it Hard?

**But networks are far more complex!**

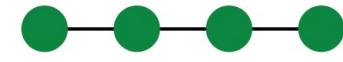
- Arbitrary size and complex topological structure (i.e., no spatial locality like grids)



Networks



Images



Text

- No fixed node ordering or reference point
- Often dynamic and have multimodal features

# Outline of Today's Lecture

**1. Basics of deep learning**



**2. Deep learning for graphs**

**3. Graph Convolutional Networks**

**4. GNNs subsume CNNs and  
Transformers**

# Stanford CS224W: Basics of Deep Learning

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>



# Machine Learning as Optimization

监督学习

特征

- **Supervised learning:** we are given input  $x$ ,  
and the goal is to predict label  $y$ .
- **Input  $x$  can be:**
  - Vectors of real numbers 向量
  - Sequences (natural language) 文本序列
  - Matrices (images) 栅格图片
  - Graphs (potentially with node and edge features) 图
- **We formulate the task as an optimization problem.**

# Machine Learning as Optimization

- Formulate the task as an optimization problem:

$$\min_{\Theta} \mathcal{L}(y, f(x))$$

Objective function

- $\Theta$ : a set of **parameters** we optimize
  - Could contain one or more scalars, vectors, matrices ...
  - E.g.  $\Theta = \{Z\}$  in the shallow encoder (the embedding lookup)

- $\mathcal{L}$ : **loss function**. Example: L2 loss

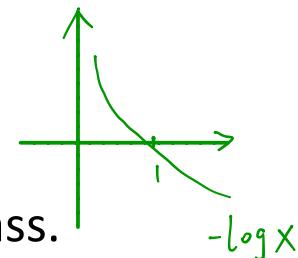
$$\mathcal{L}(y, f(x)) = \|y - f(x)\|_2$$

- Other common loss functions:
  - L1 loss, huber loss, max margin (hinge loss), cross entropy ...
  - See <https://pytorch.org/docs/stable/nn.html#loss-functions>

# Loss Function Example

多分类 交叉熵损失函数

- One common loss for classification: **cross entropy (CE)**
- Label  $y$  is a categorical vector (**one-hot encoding**)
  - e.g.  $y = \begin{array}{c|c|c|c|c} & o & o & 1 & o & o \end{array}$   $y$  is of class "3"
- $f(x) = \text{Softmax}(g(x))$ 
  - Recall from lecture 3:  $f(x)_i = \frac{e^{g(x)_i}}{\sum_{j=1}^C e^{g(x)_j}}$   $g(x)_i$  denotes  $i$ -th coordinate of the vector output of func.  $g(x)$
  - where  $C$  is the number of classes.
  - e.g.  $f(x) = \begin{array}{c|c|c|c|c} & 0.1 & 0.3 & 0.4 & 0.1 & 0.1 \end{array}$
- $\text{CE}(y, f(x)) = -\sum_{i=1}^C (y_i \log f(x)_i)$  -  $\log 0.4$ 
  - $y_i$  and  $f(x)_i$  are the **actual** and **predicted** values of the  $i$ -th class.
  - **Intuition:** the lower the loss, the closer the prediction is to one-hot
- **Total loss over all training examples:**
  - $\mathcal{L} = \sum_{(x,y) \in \mathcal{T}} \text{CE}(y, f(x))$ 
    - $\mathcal{T}$ : training set containing all pairs of data and labels  $(x, y)$



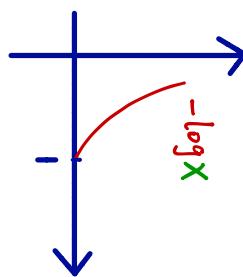
# 交叉熵损失函数

同济子豪兄 2023-3-1  
公众号 人工智能小技巧 回复 交叉熵 下载原图

$$\text{交叉熵} = -y_i \log p_i$$

|         | 猫       | 狗   | 马   | 猪   |
|---------|---------|-----|-----|-----|
| 训练集样本 1 | 标签<br>0 | 0   | 1   | 0   |
| 预测      | 0.2     | 0.6 | 0.1 | 0.1 |

$$-\log 0.6$$



$p_i$  越接近 1  
交叉熵越小

|         | 猫       | 狗   | 马   | 猪   |
|---------|---------|-----|-----|-----|
| 训练集样本 2 | 标签<br>0 | 0   | 1   | 0   |
| 预测      | 0.1     | 0.5 | 0.3 | 0.1 |

$$-\log 0.3$$

$p_i$  越接近 1  
交叉熵越小

样本之间 独立同分布

随机事件“所有样本都被正确预测”发生的概率  $P = 0.6 \times 0.3 \times 0.7$

$P$  越大  $\Rightarrow$  预测结果越好  $\Rightarrow$  算法模型越好

优化目标：更新算法模型参数 使得  $P$  最大化 极大似然估计

最大化  $P \Leftrightarrow$  最大化  $P$  取对数  $\log P = \log 0.6 + \log 0.3 + \log 0.7 \Leftrightarrow$  最小化 负对数  $-\log 0.6 - \log 0.3 - \log 0.7$

交叉熵损失函数

$$L = -\sum_i y_i \log p_i$$

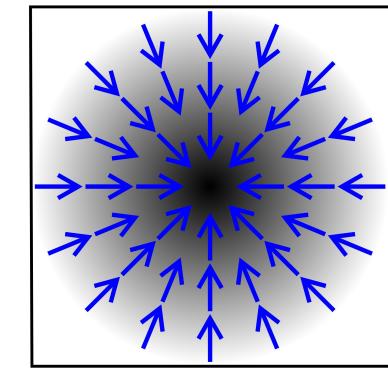
# Machine Learning as Optimization

- How to optimize the **objective function**?
- **Gradient vector:** Direction and rate of fastest increase

目标函数

Partial derivative  
偏导数

$$\nabla_{\Theta} \mathcal{L} = \left( \frac{\partial \mathcal{L}}{\partial \Theta_1}, \frac{\partial \mathcal{L}}{\partial \Theta_2}, \dots \right)$$



<https://en.wikipedia.org/wiki/Gradient>

- $\Theta_1, \Theta_2 \dots$  : components of  $\Theta$
- Recall **directional derivative**  
of a multi-variable function (e.g.  $\mathcal{L}$ ) along a given vector represents the instantaneous rate of change of the function along the vector. (斜率)
- Gradient is the directional derivative in the **direction of largest increase.** 斜率最大的方向

# Gradient Descent

梯度下降

- **Iterative algorithm:** repeatedly update weights in the (opposite) direction of gradients until convergence

$$\Theta \leftarrow \Theta - \eta \nabla_{\Theta} \mathcal{L}$$

- **Training:** Optimize  $\Theta$  iteratively
  - **Iteration:** 1 step of gradient descent
- **Learning rate (LR)  $\eta$ :** 学习率
  - Hyperparameter that controls the size of gradient step
  - Can vary over the course of training (LR scheduling)
- **Ideal termination condition:** gradient =  $\mathbf{0}$  (碗底)
  - In practice, we stop training if it no longer improves performance on **validation set** (part of dataset we hold out from training).

# Stochastic Gradient Descent (SGD)

## ■ Problem with gradient descent:

- Exact gradient requires computing  $\nabla_{\Theta} \mathcal{L}(y, f(\mathbf{x}))$ , where  $\mathbf{x}$  is the **entire** dataset! 每迭代一次，需输入所有样本计算损失函数
  - This means summing gradient contributions over all the points in the dataset
  - Modern datasets often contain billions of data points
  - Extremely expensive for every gradient descent step

## ■ Solution: Stochastic gradient descent (SGD)

- At every step, pick a different **minibatch**  $\mathcal{B}$  containing a subset of the dataset, use it as input  $\mathbf{x}$

每迭代一次，只输入batch size个样本计算损失函数

# Minibatch SGD

## ■ Concepts:

-一次迭代输入的样本数

- **Batch size**: the number of data points in a minibatch
  - E.g. number of nodes for node classification task
- **Iteration**: 1 step of SGD on a minibatch
- **Epoch**: one full pass over the dataset (# iterations is equal to ratio of dataset size and batch size)

遍历训练集全部样本

无偏估计

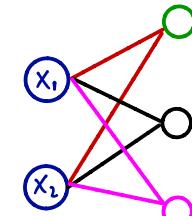
## ■ SGD is unbiased estimator of full gradient:

- But there is no guarantee on the rate of convergence
- In practice often requires tuning of learning rate
- Common optimizer that improves over SGD:
  - Adam, Adagrad, Adadelta, RMSprop ...

# Neural Network Function

- **Objective:**  $\min_{\Theta} \mathcal{L}(y, f(\mathbf{x}))$
- In deep learning, function  $f$  can be very complex
- **Example:**
  - To start simple, consider linear function  
$$f(\mathbf{x}) = \mathbf{W} \cdot \mathbf{x}, \quad \Theta = \{\mathbf{W}\}$$
  - Then, if  $f$  returns a scalar, then  $\mathbf{W}$  is a learnable **vector**  
$$\nabla_{\mathbf{W}} f = \left( \frac{\partial f}{\partial w_1}, \frac{\partial f}{\partial w_2}, \frac{\partial f}{\partial w_3} \dots \right)$$
 偏导数向量
  - But, if  $f$  returns a vector, then  $\mathbf{W}$  is the **weight matrix**

$$\nabla_{\mathbf{W}} f = \begin{bmatrix} \frac{\partial f_1}{\partial w_{11}} & \frac{\partial f_2}{\partial w_{12}} \\ \frac{\partial f_1}{\partial w_{21}} & \frac{\partial f_2}{\partial w_{22}} \end{bmatrix}$$



雅克比矩阵  
**Jacobian**

**matrix of**  $f$

# Intuition: Back Propagation

- **Goal:**  $\min_{\Theta} \mathcal{L}(y, f(\mathbf{x}))$ 
  - To minimize  $\mathcal{L}$ , we need to evaluate the gradient:  
$$\nabla_{\mathbf{W}} \mathcal{L} = \left( \frac{\partial \mathcal{L}}{\partial \mathbf{w}_1}, \frac{\partial \mathcal{L}}{\partial \mathbf{w}_2}, \frac{\partial \mathcal{L}}{\partial \mathbf{w}_3}, \dots \right)$$
which means we need to derive derivative of  $\mathcal{L}$ .
- **Overview of Back-propagation:** 复合函数求偏导
  - $\mathcal{L}$  is composed from some set of predefined building block functions  $g(\cdot)$
  - For each such  $g$  we also have its derivative  $g'$
  - Then we can automatically compute  $\nabla_{\Theta} \mathcal{L}$  by evaluating appropriate funcs.  $g'$  on the minibatch  $\mathcal{B}$ .

# Back-propagation

$$f = g \cdot h$$

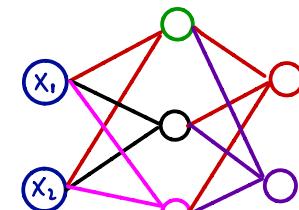
- How about a more complex function:

$$f(\mathbf{x}) = W_2(W_1 \mathbf{x}), \Theta = \{W_1, W_2\}$$

- Recall chain rule:

$$\frac{df}{dx} = \frac{dg}{dh} \cdot \frac{dh}{dx} \text{ or } f'(x) = g'(h(x))h'(x)$$

- Example:  $\nabla_{\mathbf{x}} f = \frac{\partial f}{\partial (W_1 \mathbf{x})} \cdot \frac{\partial (W_1 \mathbf{x})}{\partial \mathbf{x}}$



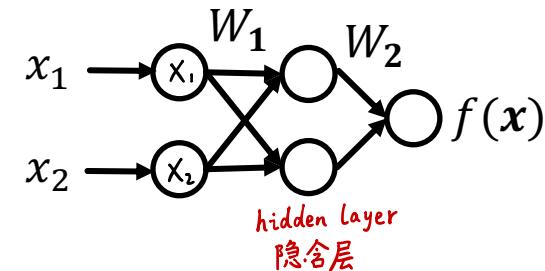
链式法则

- Back-propagation: Use of chain rule to propagate gradients of intermediate steps, and finally obtain gradient of  $\mathcal{L}$  w.r.t.  $\Theta$ .

In other words:  
 $f(\mathbf{x}) = W_2(W_1 \mathbf{x})$   
 $h(x) = W_1 \mathbf{x}$   
 $g(z) = W_2 z$

# Back-propagation Example (1)

- Example: Simple 2-layer linear network
- $f(\mathbf{x}) = g(h(\mathbf{x})) = W_2(W_1 \mathbf{x})$   
标量  $| \times |$        $1 \times 2$      $2 \times 2$      $2 \times 1$
- $\mathcal{L} = \sum_{(x,y) \in \mathcal{B}} \left\| (y, -f(x)) \right\|_2$ 
  - The loss  $\mathcal{L}$  sums the L2 loss in a minibatch  $\mathcal{B}$ .
- Hidden layer:
  - Intermediate representation of input  $\mathbf{x}$
  - Here we use  $h(\mathbf{x}) = W_1 \mathbf{x}$  to denote the hidden layer
  - $f(\mathbf{x}) = W_2 h(\mathbf{x})$



# Back-propagation Example (2)

- **Forward propagation:** 前向预测 求损失函数

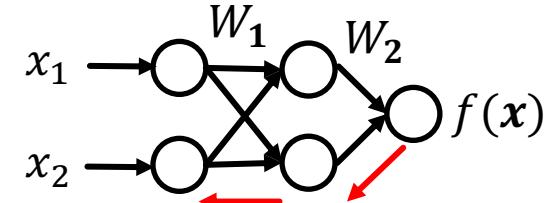
Compute loss starting from input



Remember:  
 $f(x) = W_2(W_1 x)$

$$h(x) = W_1 x$$

$$g(z) = W_2 z$$



- **Back-propagation to compute gradient of**

反向传播 求梯度

$$\Theta = \{W_1, W_2\}$$

Start from loss, compute the gradient

$$\frac{\partial \mathcal{L}}{\partial W_2} = \frac{\partial \mathcal{L}}{\partial f} \cdot \frac{\partial f}{\partial W_2},$$



Compute backwards

$$\frac{\partial \mathcal{L}}{\partial W_1} = \frac{\partial \mathcal{L}}{\partial f} \cdot \frac{\partial f}{\partial W_2} \cdot \frac{\partial W_2}{\partial W_1}$$



Compute backwards

# Non-linearity

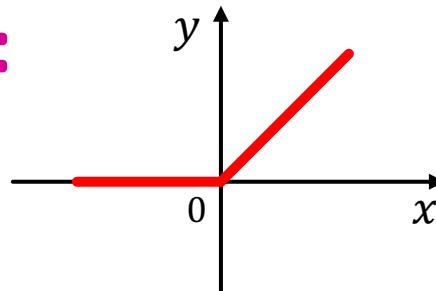
- Note that in  $f(\mathbf{x}) = W_2(W_1 \mathbf{x})$ ,  $W_2 W_1$  is another matrix (vector, if we do binary classification)
  - Hence  $f(\mathbf{x})$  is still linear w.r.t.  $\mathbf{x}$  no matter how many weight matrices we compose 矩阵代表线性变换

## 非线性激活函数

- We introduce non-linearity:

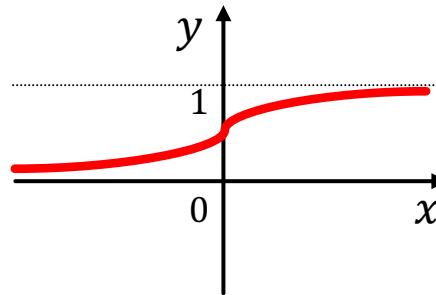
- Rectified linear unit (ReLU)

$$ReLU(x) = \max(x, 0)$$



- Sigmoid

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

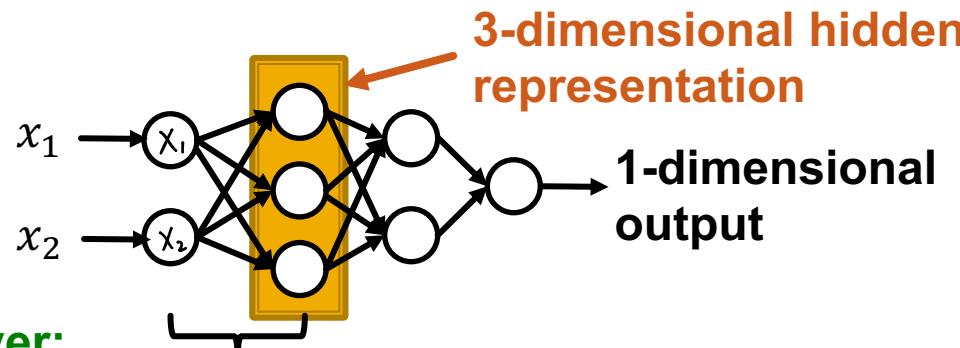


# Multi-layer Perceptron (MLP)

- Each layer of MLP combines linear transformation and non-linearity:

$$\mathbf{x}^{(l+1)} = \sigma(W_l \mathbf{x}^{(l)} + b^l)$$

- where  $W_l$  is weight matrix that transforms hidden representation at layer  $l$  to layer  $l + 1$
- $b^l$  is bias at layer  $l$ , and is added to the linear transformation of  $\mathbf{x}$  偏置项
- $\sigma$  is non-linearity function (e.g., sigmod) 非线性激活函数
- Suppose  $\mathbf{x}$  is 2-dimensional, with entries  $x_1$  and  $x_2$



Every layer:  
Linear transformation +  
non-linearity

# Summary

- **Objective function:**

$$\min_{\Theta} \mathcal{L}(y, f(\mathbf{x}))$$

- $f$  can be a simple linear layer, an MLP, or other neural networks (e.g., a GNN later)
- Sample a minibatch of input  $\mathbf{x}$
- **Forward propagation:** Compute  $\mathcal{L}$  given  $\mathbf{x}$
- **Back-propagation:** Obtain gradient  $\nabla_{\mathbf{W}} \mathcal{L}$  using a chain rule.
- Use **stochastic gradient descent (SGD)** to optimize for  $\Theta$  over many iterations.

# Outline of Today's Lecture

1. Basics of deep learning



2. Deep learning for graphs



3. Graph Convolutional Networks

4. GNNs subsume CNNs and  
Transformers

# Stanford CS224W: Deep Learning for Graphs

CS224W: Machine Learning with Graphs

Jure Leskovec, Stanford University

<http://cs224w.stanford.edu>



# Content

- **Local network neighborhoods:**
  - Describe aggregation strategies
  - Define computation graphs
- **Stacking multiple layers:**
  - Describe the model, parameters, training
  - How to fit the model?
  - Simple example for unsupervised and supervised training

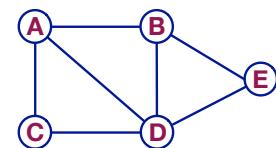
# Setup

- Assume we have a graph  $G$ :
  - $V$  is the **vertex set**
  - $A$  is the **adjacency matrix** (assume binary)
  - $X \in \mathbb{R}^{m \times |V|}$  is a matrix of **node features** 节点属性特征
  - $v$ : a node in  $V$ ;  $N(v)$ : the set of neighbors of  $v$ .
  - **Node features:** 节点  $v$  的邻域 (相连节点.)
    - Social networks: User profile, User image
    - Biological networks: Gene expression profiles, gene functional information
    - When there is no node feature in the graph dataset:
      - Indicator vectors (one-hot encoding of a node)
      - Vector of constant 1: [1, 1, ..., 1]

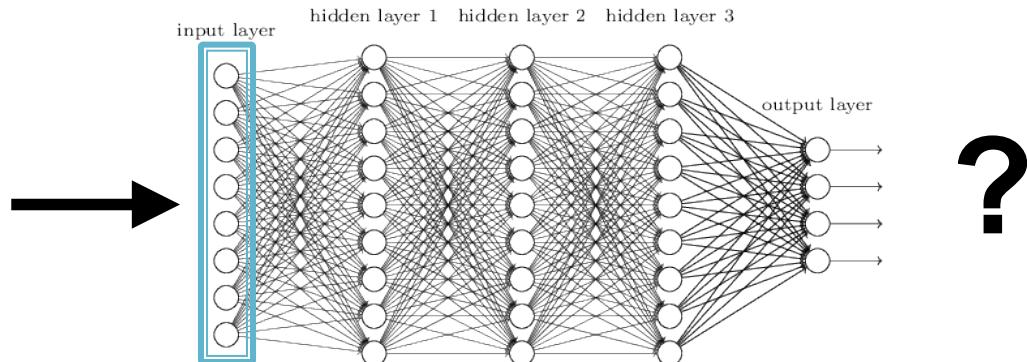
# A Naïve Approach

朴素想法：直接输入邻接矩阵A

- Join adjacency matrix and features
- Feed them into a deep neural net:



|   | A | B | C | D | E | Feat |
|---|---|---|---|---|---|------|
| A | 0 | 1 | 1 | 1 | 0 | 1 0  |
| B | 1 | 0 | 0 | 1 | 1 | 0 0  |
| C | 1 | 0 | 0 | 1 | 0 | 0 1  |
| D | 1 | 1 | 1 | 0 | 1 | 1 1  |
| E | 0 | 1 | 0 | 1 | 0 | 1 0  |

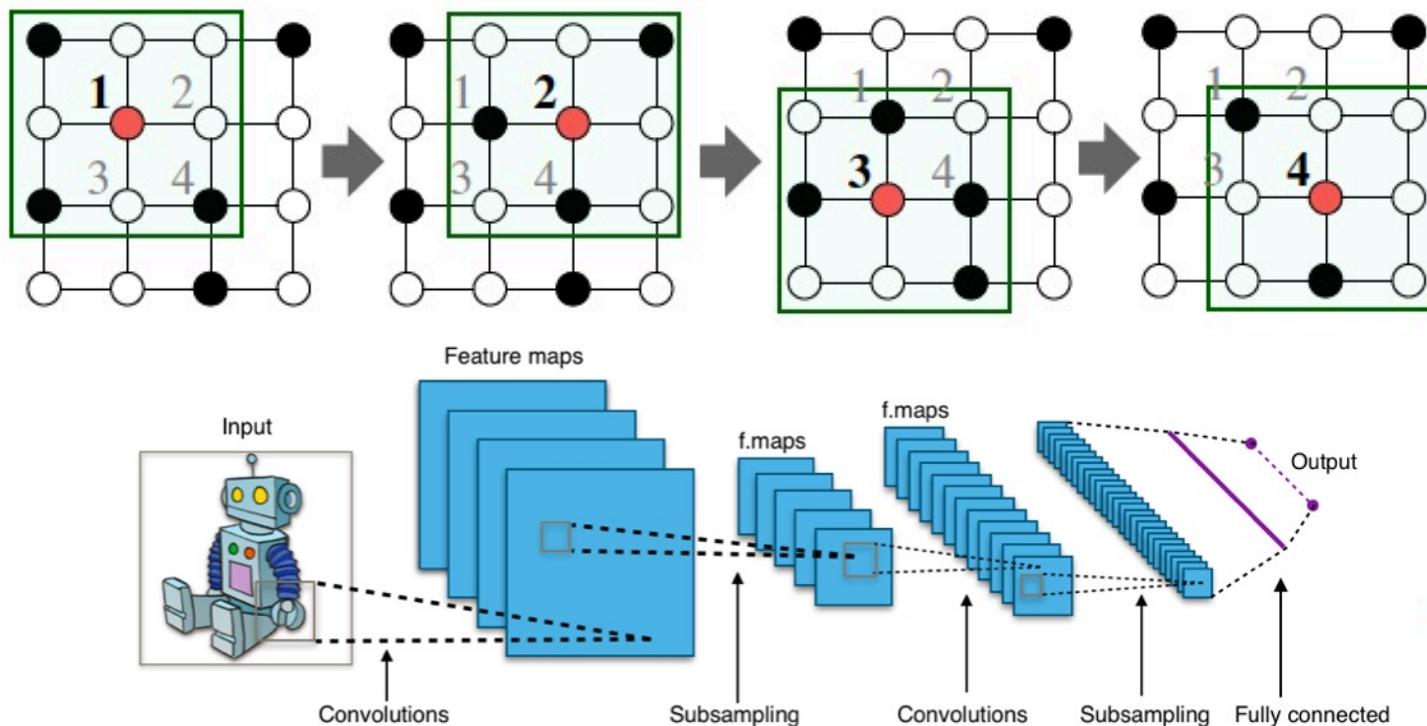


- Issues with this idea:

- $O(|V|)$  parameters      过拟合
- Not applicable to graphs of different sizes      无法泛化到新节点
- Sensitive to node ordering      不具备“变换不变性”

# Idea: Convolutional Networks

CNN on an image:

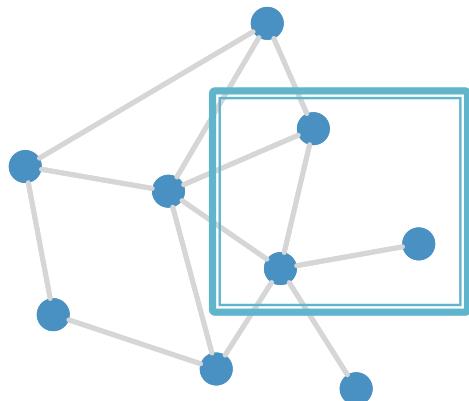


Goal is to generalize convolutions beyond simple lattices  
Leverage node features/attributes (e.g., text, images)

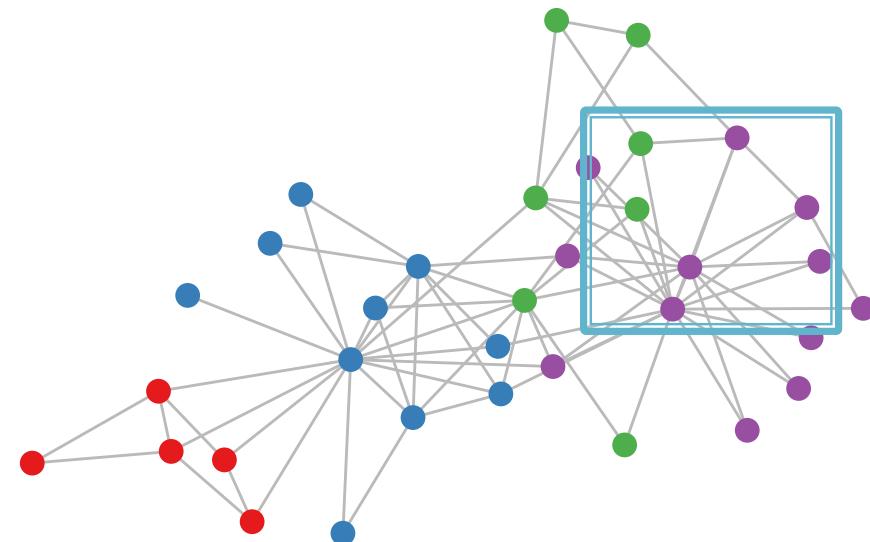
栅格

# Real-World Graphs

But our graphs look like this:



or this:



- There is no fixed notion of locality or sliding window on the graph
- Graph is permutation invariant 固定的  
图具有“变换不变性”

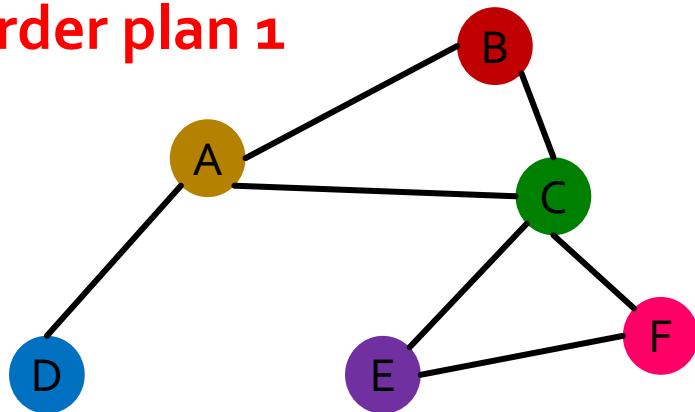
# Permutation Invariance

- Graph does not have a canonical order of the nodes!
- We can have many different order plans.

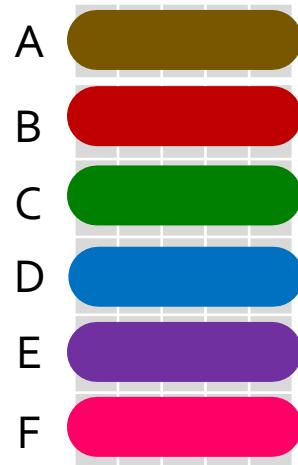
# Permutation Invariance

- Graph does not have a canonical order of the nodes!

Order plan 1



Node features  $X_1$



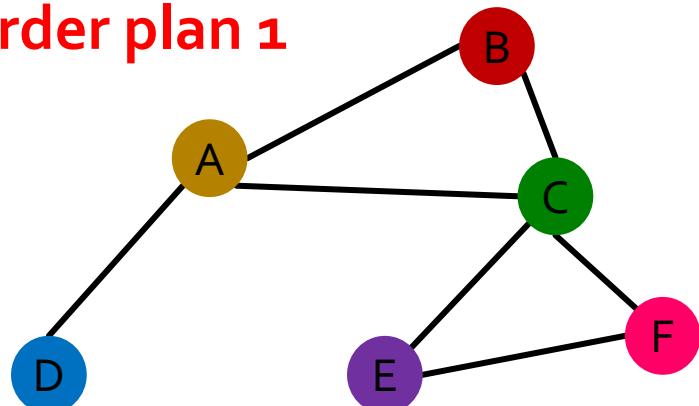
Adjacency matrix  $A_1$

|   | A    | B    | C    | D    | E    | F    |
|---|------|------|------|------|------|------|
| A | Gray | Blue | Blue | Blue | Gray | Gray |
| B | Blue | Gray | Gray | Gray | Gray | Gray |
| C | Blue | Blue | Gray | Gray | Blue | Blue |
| D | Blue | Gray | Gray | Gray | Gray | Gray |
| E | Gray | Gray | Blue | Gray | Gray | Blue |
| F | Gray | Gray | Gray | Blue | Gray | Gray |

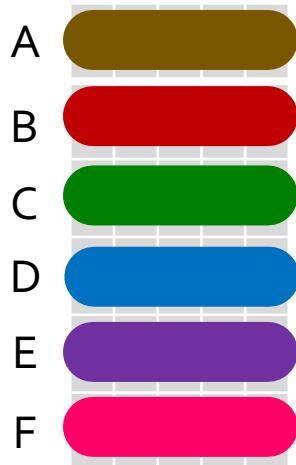
# Permutation Invariance

- Graph does not have a canonical order of the nodes!

Order plan 1



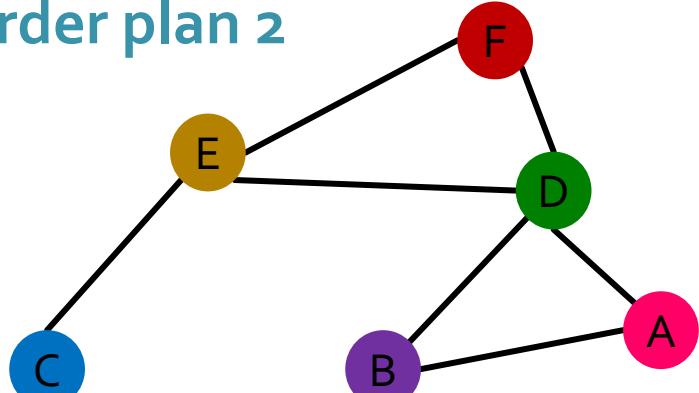
Node features  $X_1$



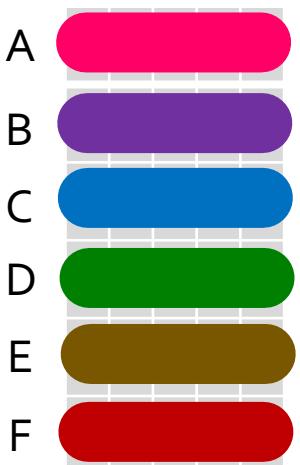
Adjacency matrix  $A_1$

|   | A    | B    | C    | D    | E    | F    |
|---|------|------|------|------|------|------|
| A | Gray | Blue | Blue | Blue | Blue | Gray |
| B | Blue | Gray | Blue | Blue | Blue | Gray |
| C | Blue | Blue | Gray | Blue | Blue | Blue |
| D | Blue | Blue | Blue | Gray | Blue | Blue |
| E | Gray | Gray | Blue | Blue | Gray | Blue |
| F | Gray | Gray | Blue | Blue | Blue | Gray |

Order plan 2



Node features  $X_2$



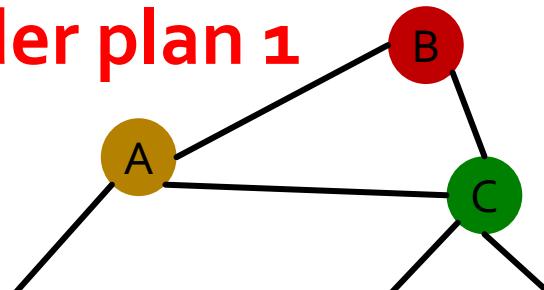
Adjacency matrix  $A_2$

|   | A    | B    | C    | D    | E    | F    |
|---|------|------|------|------|------|------|
| A | Gray | Blue | Gray | Blue | Gray | Gray |
| B | Blue | Gray | Gray | Gray | Blue | Gray |
| C | Gray | Gray | Gray | Gray | Gray | Blue |
| D | Blue | Blue | Gray | Gray | Blue | Blue |
| E | Gray | Gray | Blue | Blue | Gray | Blue |
| F | Gray | Gray | Blue | Blue | Blue | Gray |

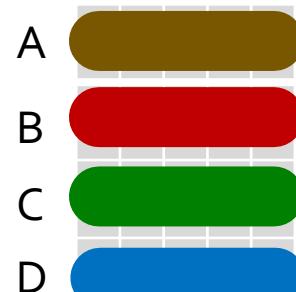
# Permutation Invariance

- Graph does not have a canonical order of the nodes!

Order plan 1



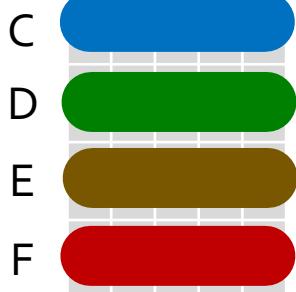
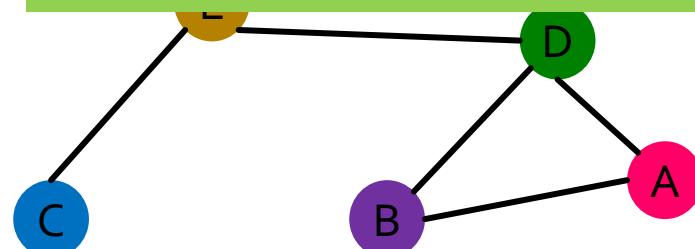
Node feature  $X_1$



Adjacency matrix  $A_1$

|   | A    | B    | C    | D    | E    | F    |
|---|------|------|------|------|------|------|
| A | Gray | Blue | Blue | Blue | Blue | Gray |
| B | Blue | Gray | Blue | Blue | Blue | Gray |
| C | Blue | Blue | Gray | Blue | Blue | Gray |
| D | Blue | Gray | Gray | Gray | Blue | Gray |

Graph and node representations  
should be the same for Order plan 1  
and Order plan 2



|   | B    | C    | D    | E    | F    |
|---|------|------|------|------|------|
| B | Blue | Gray | Gray | Blue | Gray |
| C | Gray | Gray | Gray | Blue | Gray |
| D | Blue | Blue | Gray | Gray | Blue |
| E | Gray | Gray | Blue | Blue | Gray |
| F | Gray | Gray | Gray | Blue | Gray |

# Permutation Invariance

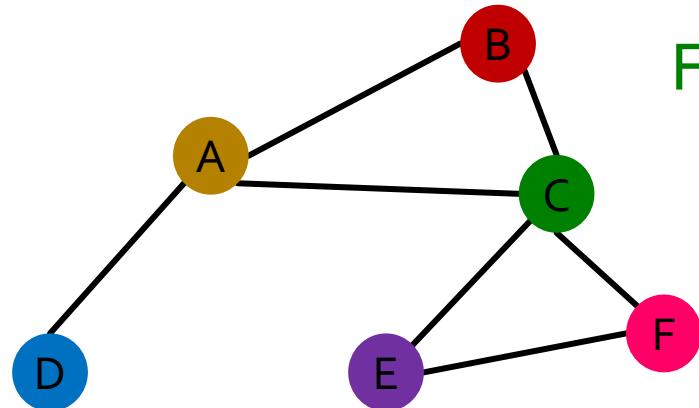
What does it mean by “graph representation is same for two order plans”?

- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a vector  $\mathbb{R}^d$  then

$$f(A_1, X_1) = f(A_2, X_2)$$

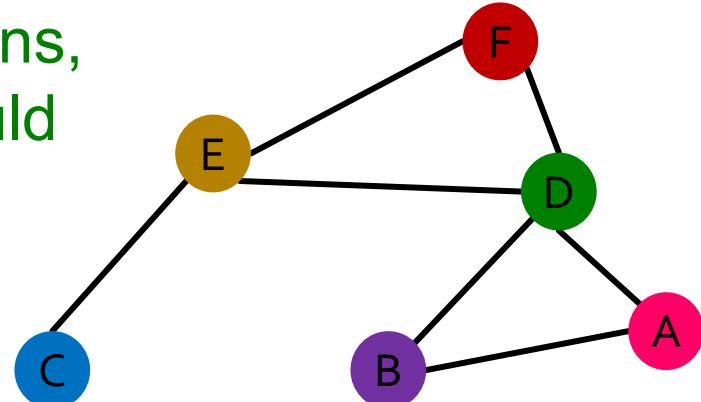
$A$  is the adjacency matrix  
 $X$  is the node feature matrix

Order plan 1:  $A_1, X_1$



For two order plans,  
output of  $f$  should  
be the same!  
与节点编号顺序无关  
与图的显示方式无关

Order plan 2:  $A_2, X_2$



# Permutation Invariance

What does it mean by “graph representation is same for two order plans”?

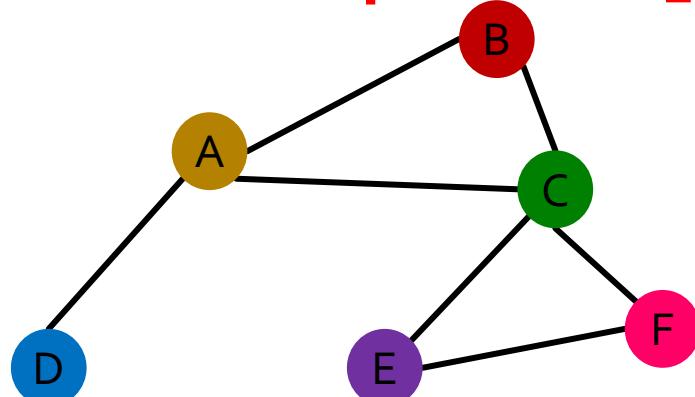
- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a vector  $\mathbb{R}^d$ .  
 $A$  is the adjacency matrix  
 $X$  is the node feature matrix
- Then, if  $f(A_i, X_i) = f(A_j, X_j)$  for any order plan  $i$  and  $j$ , we formally say  $f$  is a **permutation invariant function**.

For a graph with  $m$  nodes, there are  $m!$  different order plans.

# Permutation Equivariance

**Similarly for node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

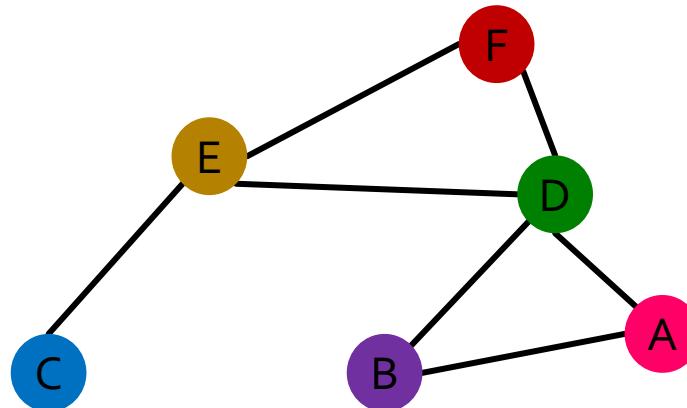
Order plan 1:  $A_1, X_1$



$$f(A_1, X_1) =$$

|   |  |  |
|---|--|--|
| A |  |  |
| B |  |  |
| C |  |  |
| D |  |  |
| E |  |  |
| F |  |  |

Order plan 2:  $A_2, X_2$



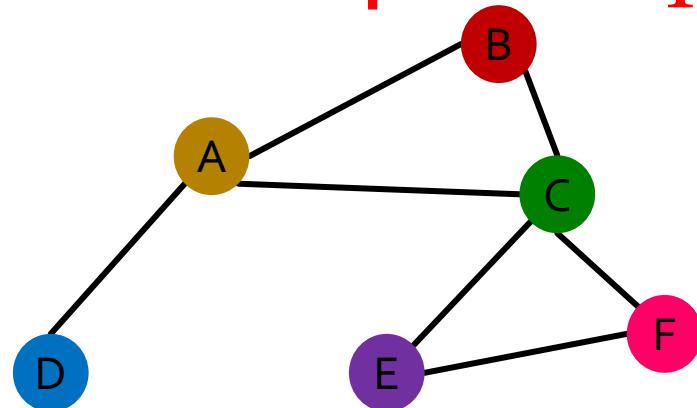
$$f(A_2, X_2) =$$

|   |  |  |
|---|--|--|
| A |  |  |
| B |  |  |
| C |  |  |
| D |  |  |
| E |  |  |
| F |  |  |

# Permutation Equivariance

**Similarly for node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

Order plan 1:  $A_1, X_1$

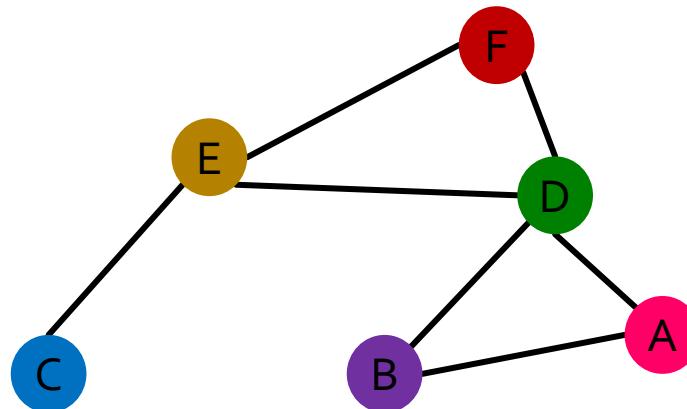


Representation vector  
of the brown node A

|   |          |          |
|---|----------|----------|
| A | [brown]  | [brown]  |
| B | [red]    | [red]    |
| C | [green]  | [green]  |
| D | [blue]   | [blue]   |
| E | [purple] | [purple] |
| F | [pink]   | [pink]   |

$$f(A_1, X_1) =$$

Order plan 2:  $A_2, X_2$



C

E

B

D

A

F

$$f(A_2, X_2) =$$

|   |          |          |
|---|----------|----------|
| A | [red]    | [red]    |
| B | [purple] | [purple] |
| C | [blue]   | [blue]   |
| D | [green]  | [green]  |
| E | [brown]  | [brown]  |
| F | [pink]   | [pink]   |

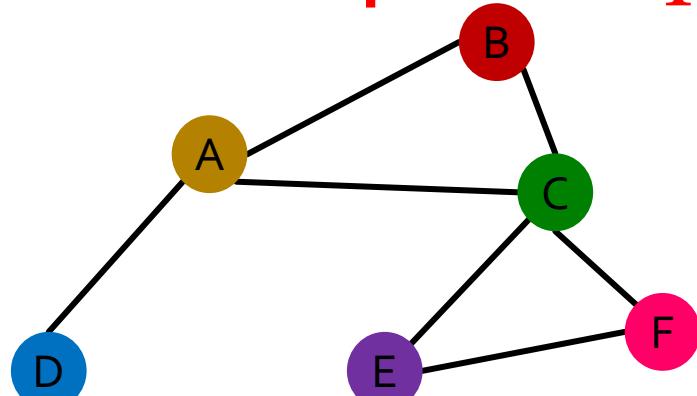
Representation vector  
of the brown node E

For two order plans, the vector of node  
at the same position is the same!

# Permutation Equivariance

**Similarly for node representation:** We learn a function  $f$  that maps nodes of  $G$  to a matrix  $\mathbb{R}^{m \times d}$ .

Order plan 1:  $A_1, X_1$

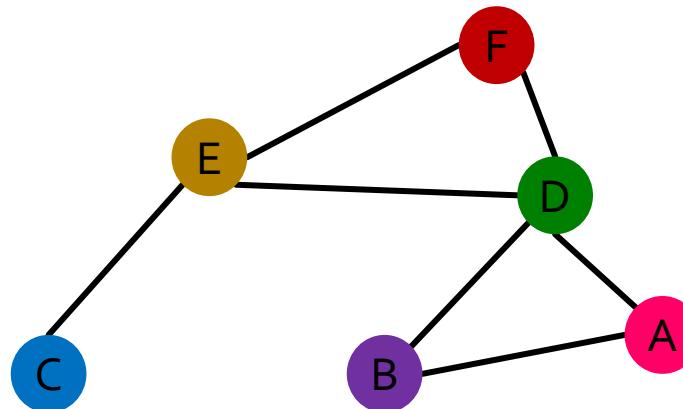


$$f(A_1, X_1) = \begin{matrix} & \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} \\ \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & \\ \boxed{\text{B}} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & \\ \text{C} & \text{C} & \text{D} & \text{E} & \text{F} & & \\ \text{D} & \text{D} & \text{E} & \text{F} & & & \\ \text{E} & \text{E} & \text{F} & & & & \\ \text{F} & \text{F} & & & & & \end{matrix}$$

Representation vector of the brown node C

For two order plans, the vector of node at the same position is the same!

Order plan 2:  $A_2, X_2$



$$f(A_2, X_2) = \begin{matrix} & \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} \\ \text{A} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & \\ \boxed{\text{D}} & \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & \\ \text{B} & \text{C} & \text{D} & \text{E} & \text{F} & & \\ \text{C} & \text{D} & \text{E} & \text{F} & & & \\ \text{D} & \text{E} & \text{F} & & & & \\ \text{E} & \text{F} & & & & & \\ \text{F} & & & & & & \end{matrix}$$

Representation vector of the brown node D

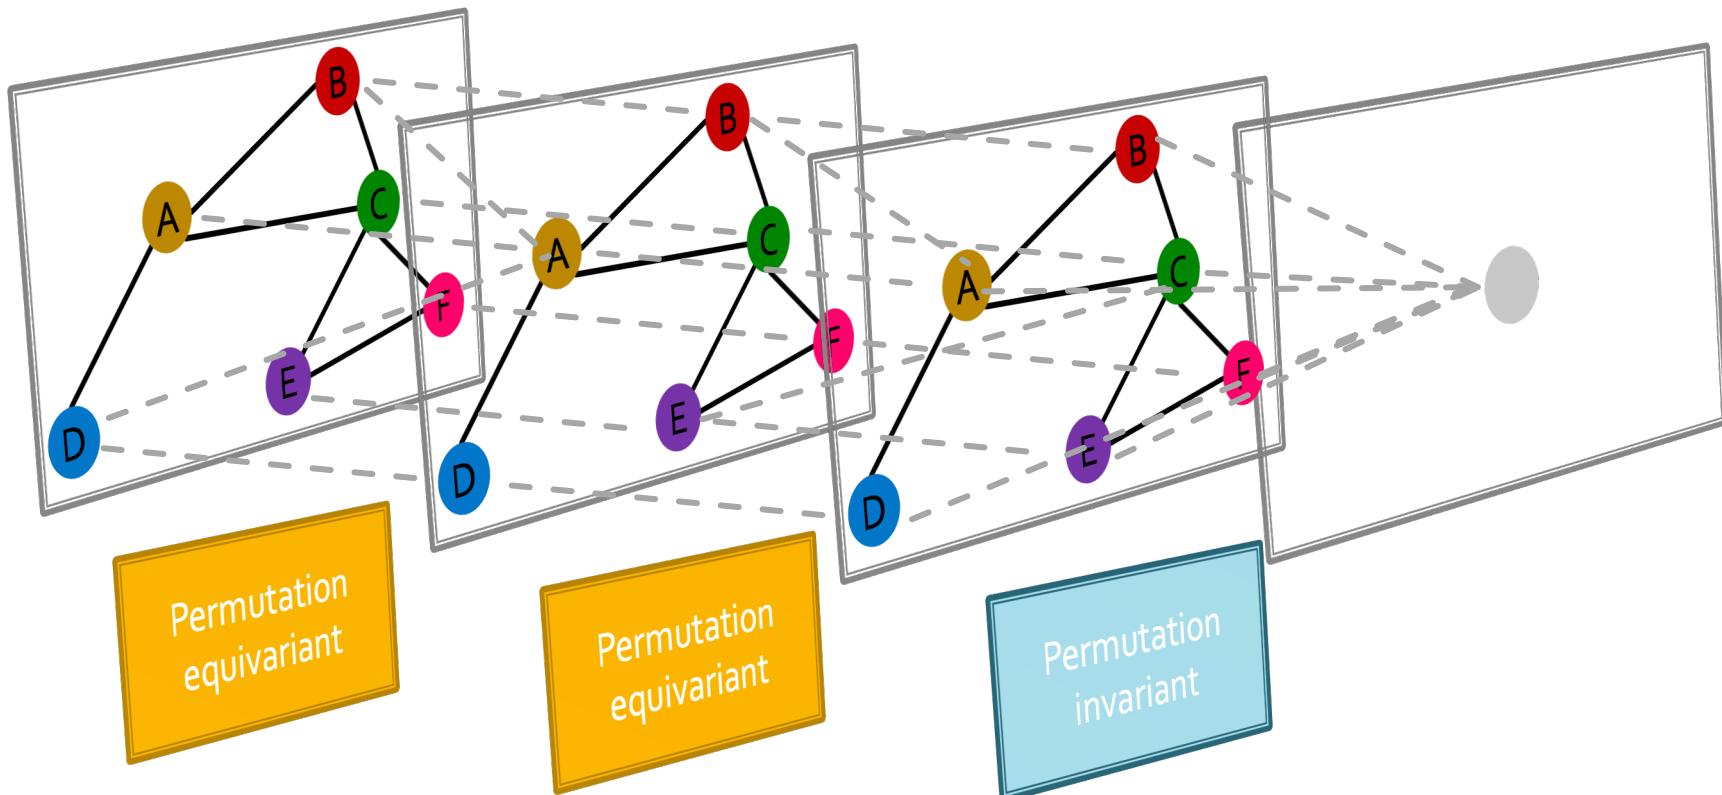
# Permutation Equivariance

## For node representation

- Consider we learn a function  $f$  that maps a graph  $G = (A, X)$  to a matrix  $\mathbb{R}^{m \times d}$ 
  - graph has  $m$  nodes, each row is the embedding of a node.
- Similarly, if this property holds for any pair of order plan  $i$  and  $j$ , we say  $f$  is a **permutation equivariant function**.

# Graph Neural Network Overview

- Graph neural networks consist of multiple permutation equivariant / invariant functions.

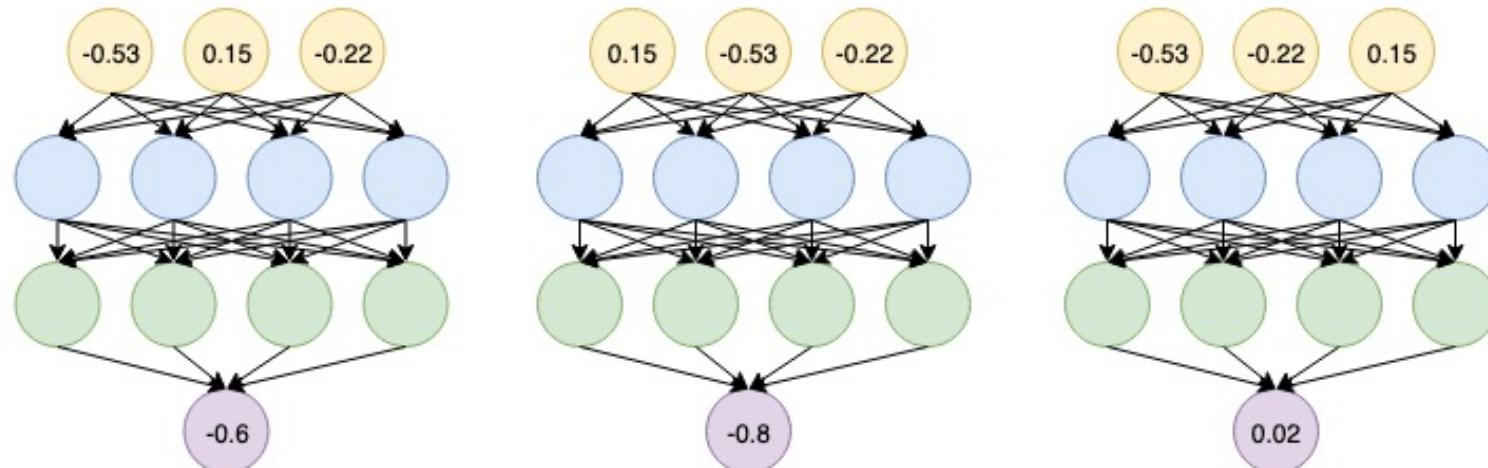


# Graph Neural Network Overview

Are other neural network architectures, e.g.,  
MLPs, permutation invariant / equivariant?

- No.

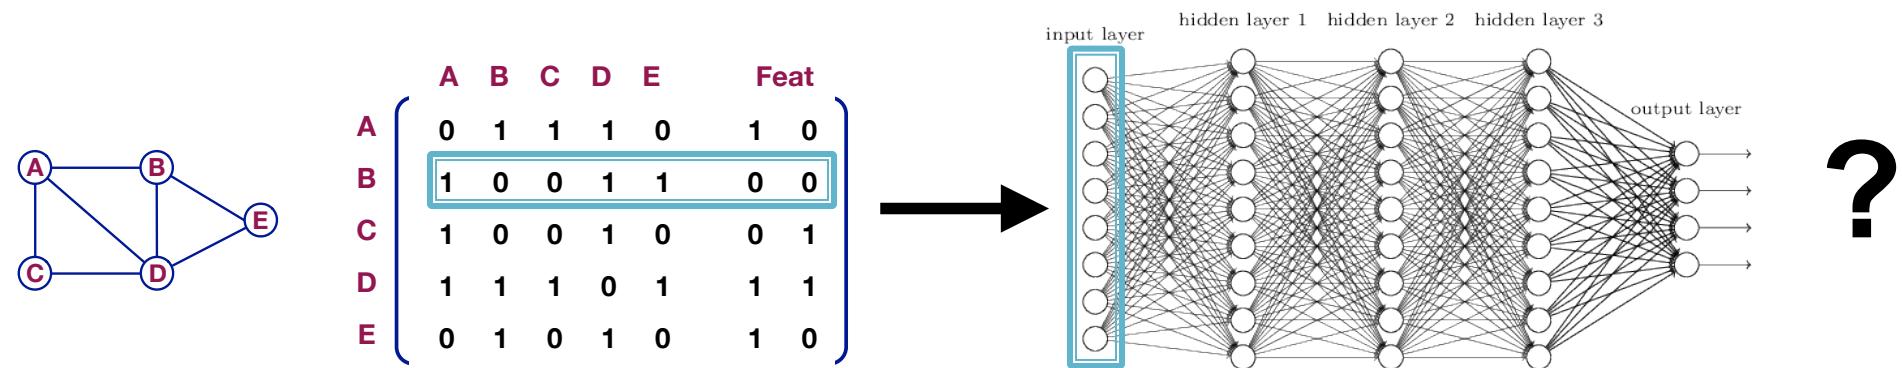
Switching the order of the  
input leads to different  
outputs!



# Graph Neural Network Overview

Are other neural network architectures, e.g.,  
MLPs, permutation invariant / equivariant?

- No.



This explains why **the naïve MLP approach fails for graphs!**

# Graph Neural Network Overview

- Are any neural network architecture, e.g.,

Next: Design graph neural networks that are permutation invariant / equivariant by **passing and aggregating information from neighbors!**

消息传递和聚合

?

# Outline of Today's Lecture

1. Basics of deep learning



2. Deep learning for graphs



图卷积神经网络

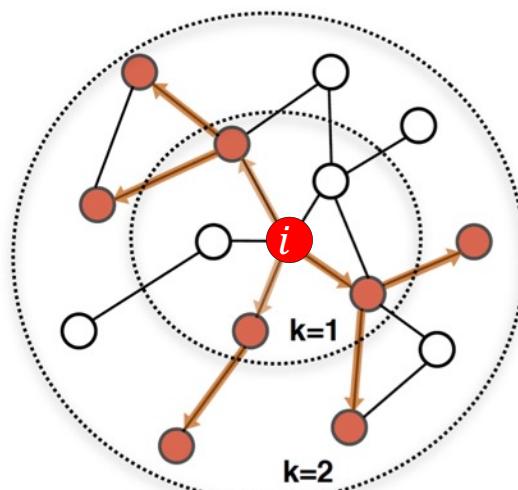
3. Graph Convolutional Networks



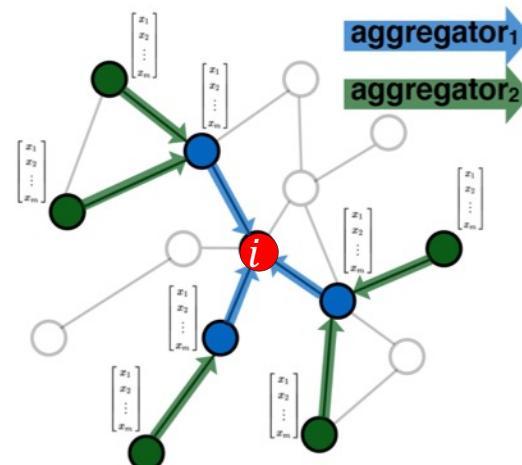
4. GNNs subsume CNNs and  
Transformers

# Graph Convolutional Networks

Idea: Node's neighborhood defines a computation graph 计算图



Determine node computation graph

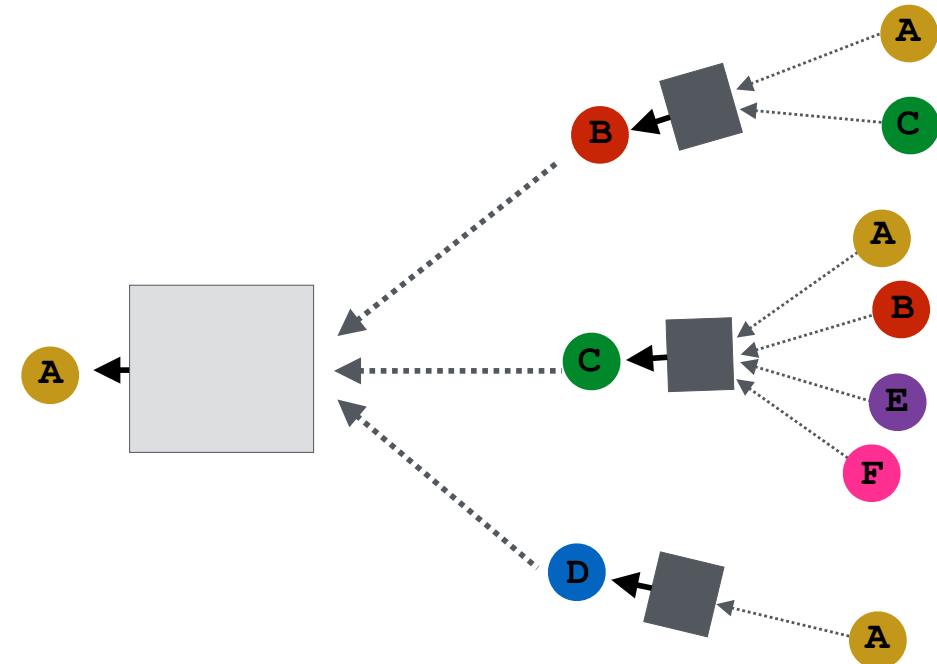
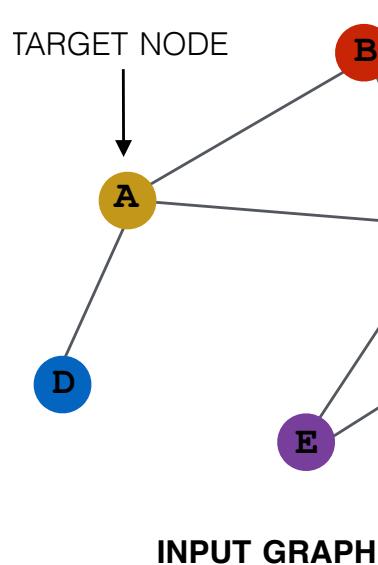


Propagate and transform information

Learn how to propagate information across the graph to compute node features

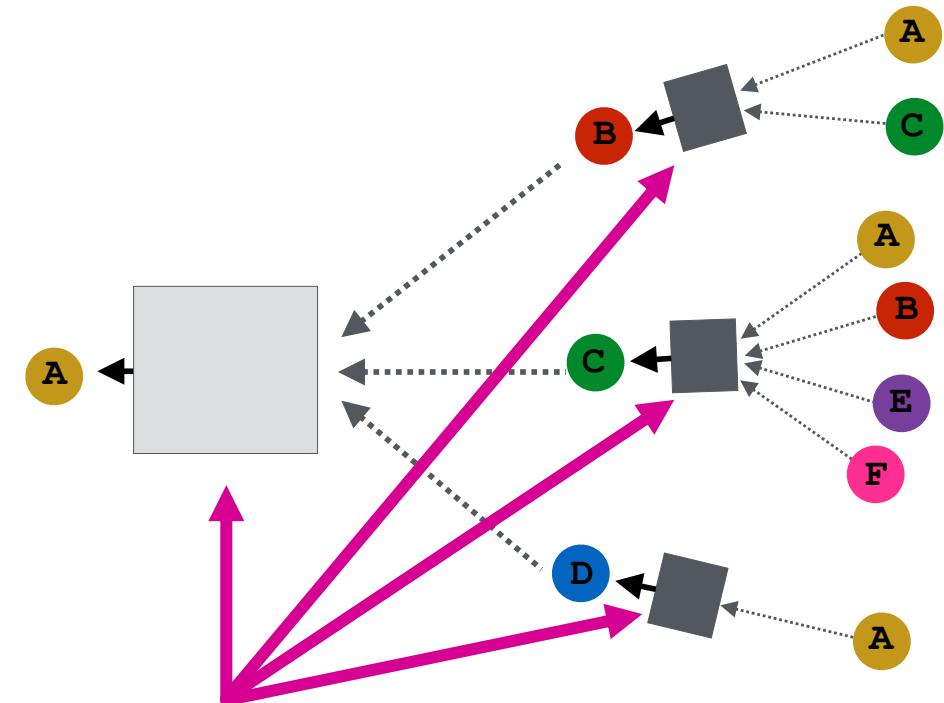
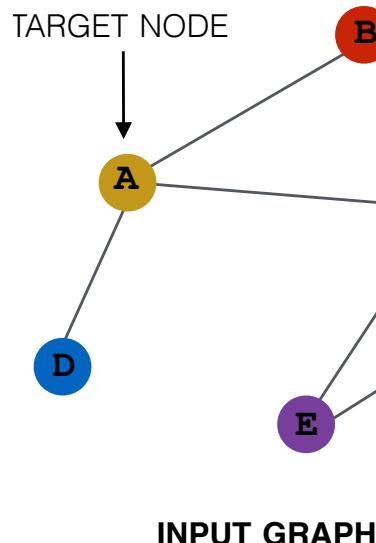
# Idea: Aggregate Neighbors

- **Key idea:** Generate node embeddings based on **local network neighborhoods**



# Idea: Aggregate Neighbors

- **Intuition:** Nodes aggregate information from their neighbors using neural networks



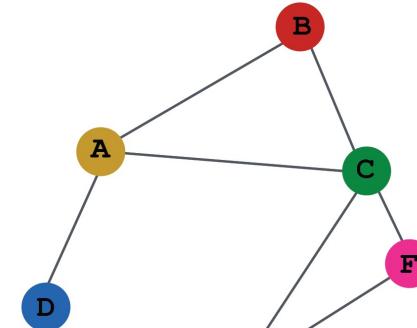
**Neural networks**

# Idea: Aggregate Neighbors

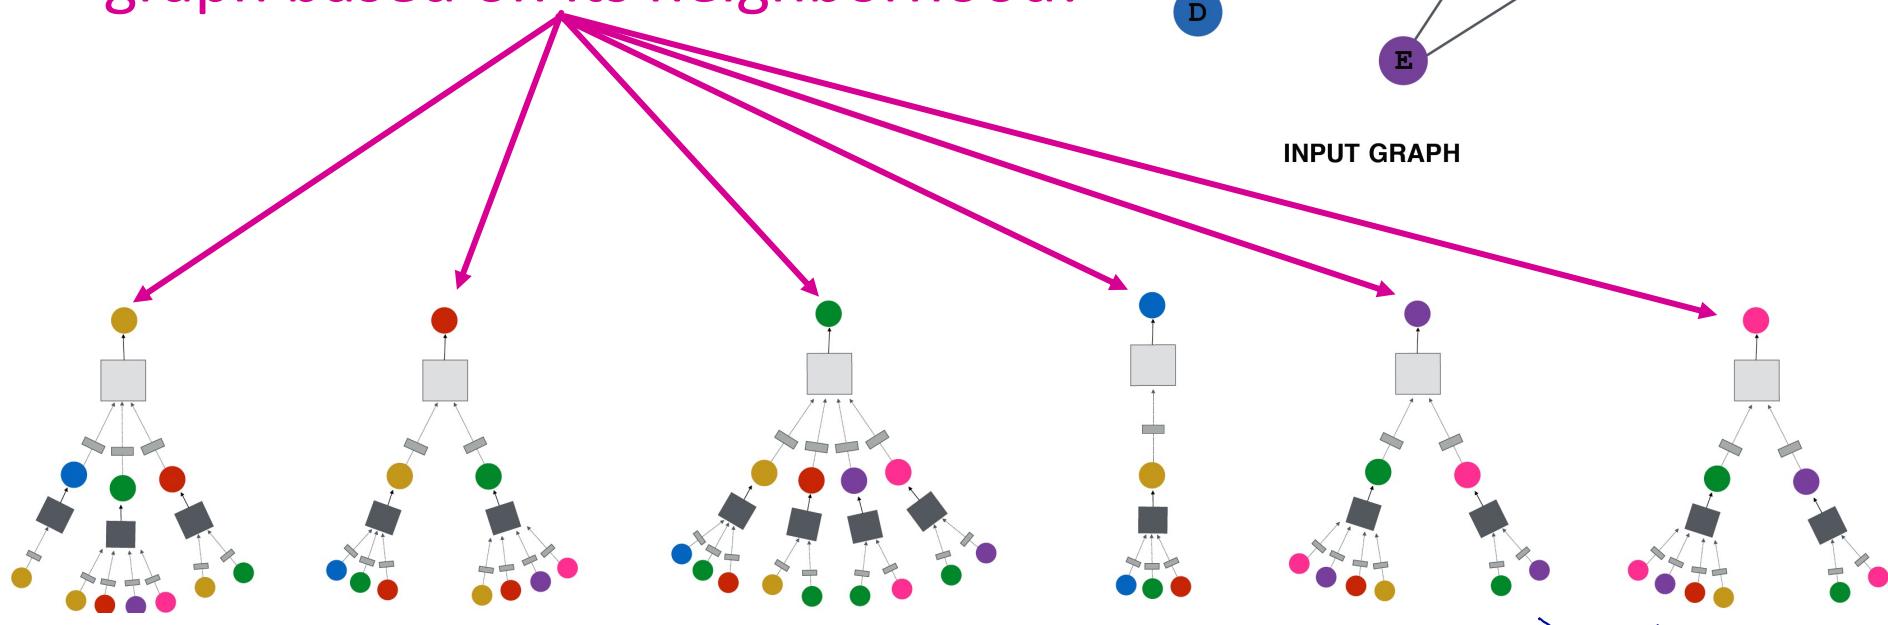
- **Intuition:** Network neighborhood defines a computation graph

每个节点分别构建自己的计算图

Every node defines a computation graph based on its neighborhood!

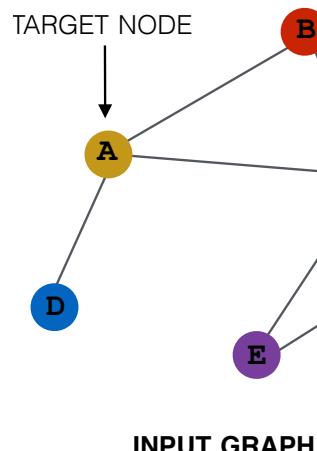


INPUT GRAPH

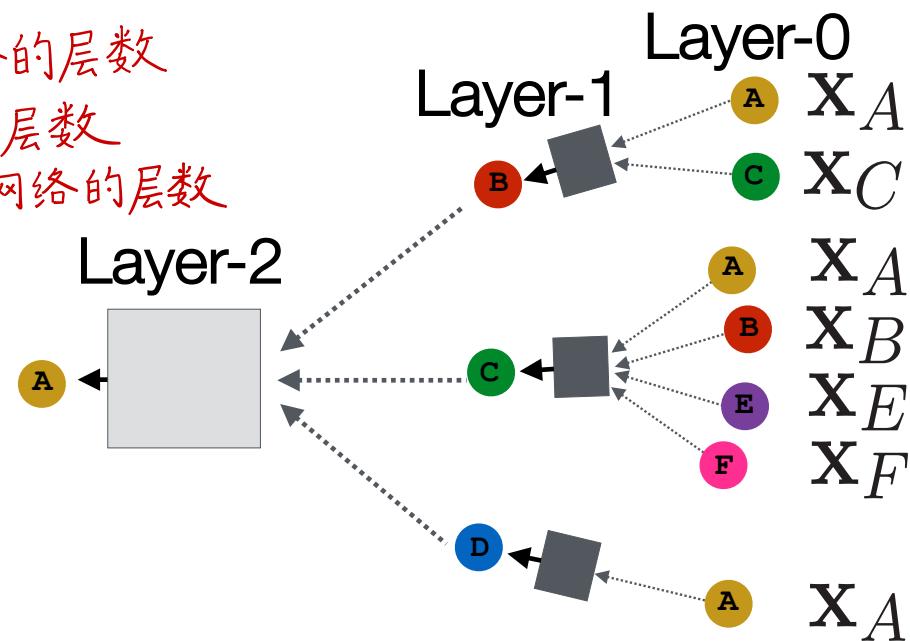


# Deep Model: Many Layers

- Model can be of arbitrary depth:
  - Nodes have embeddings at each layer
  - Layer-0 embedding of node  $v$  is its input feature,  $x_v$
  - Layer- $k$  embedding gets information from nodes that are  $k$  hops away

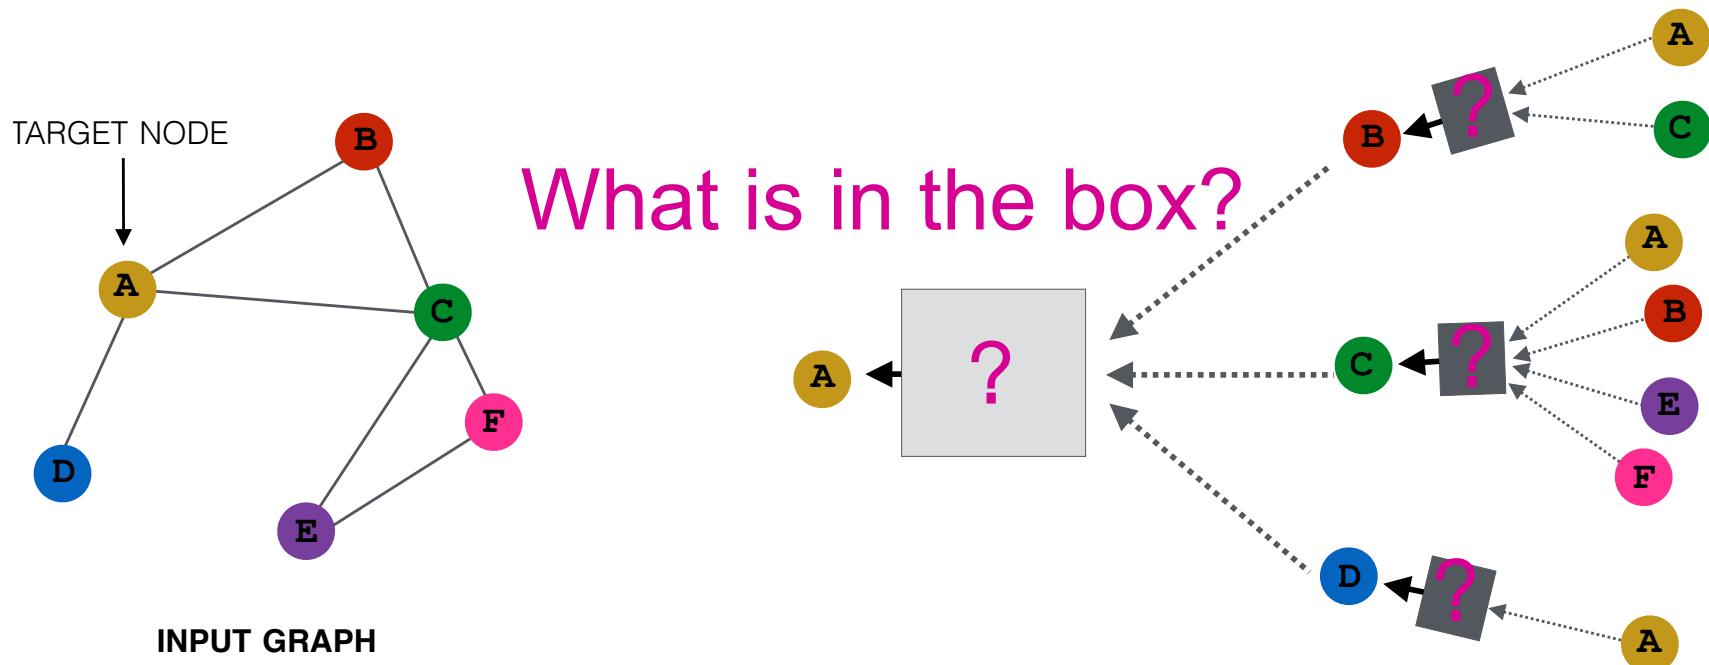


图神经网络的层数  
是计算图的层数  
而不是神经网络的层数



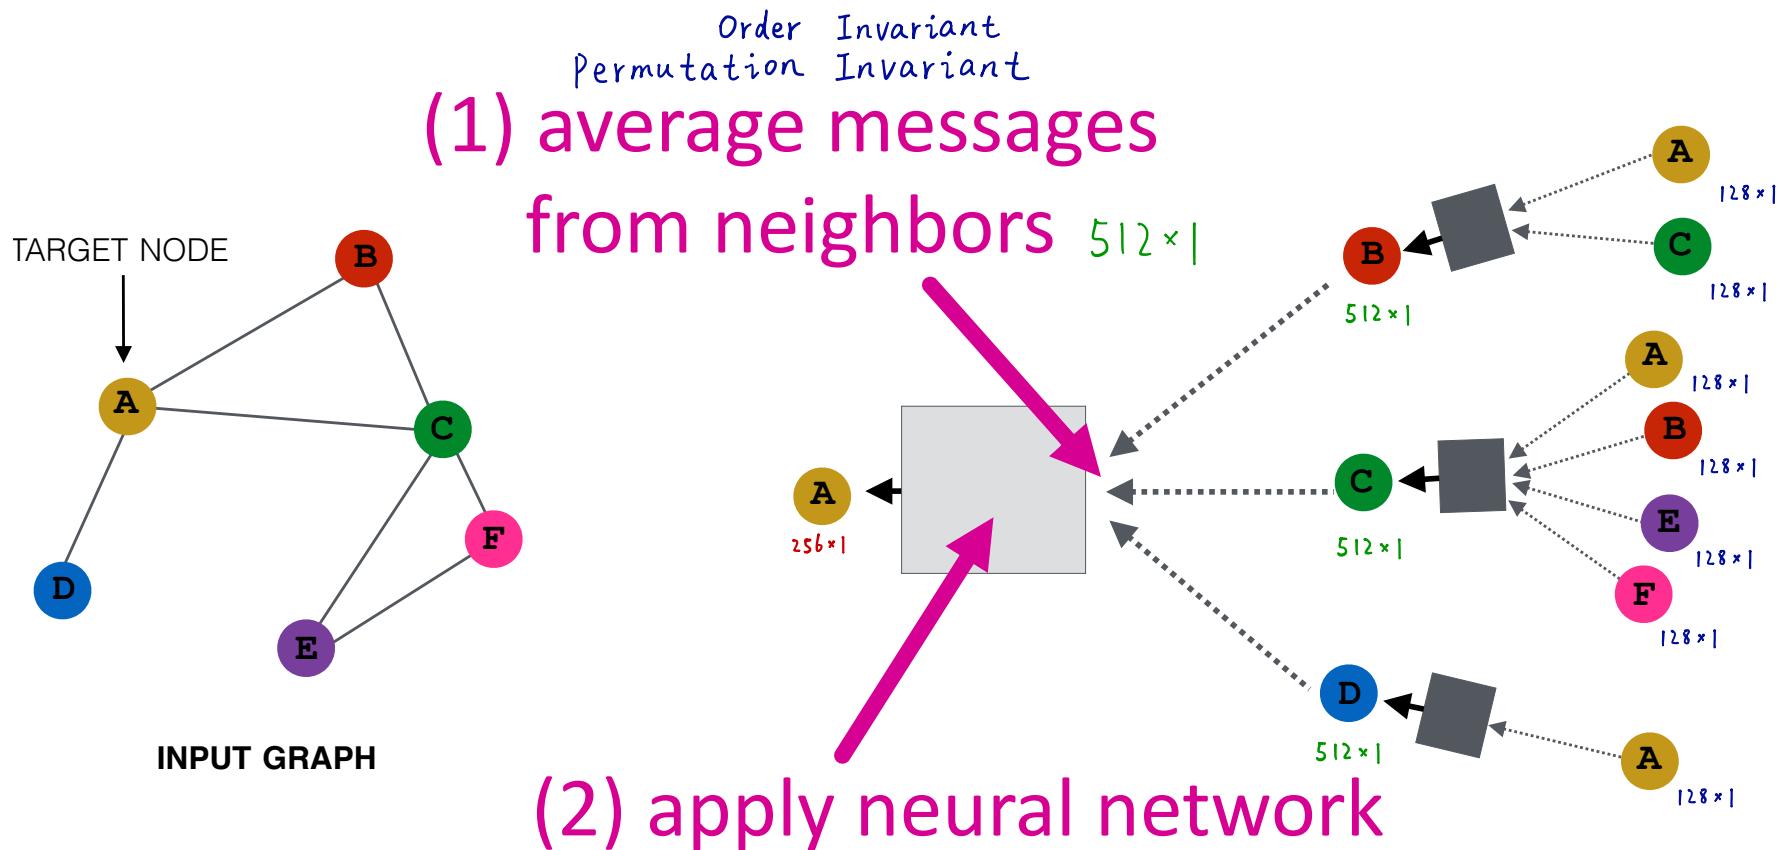
# Neighborhood Aggregation

- **Neighborhood aggregation:** Key distinctions are in how different approaches aggregate information across the layers



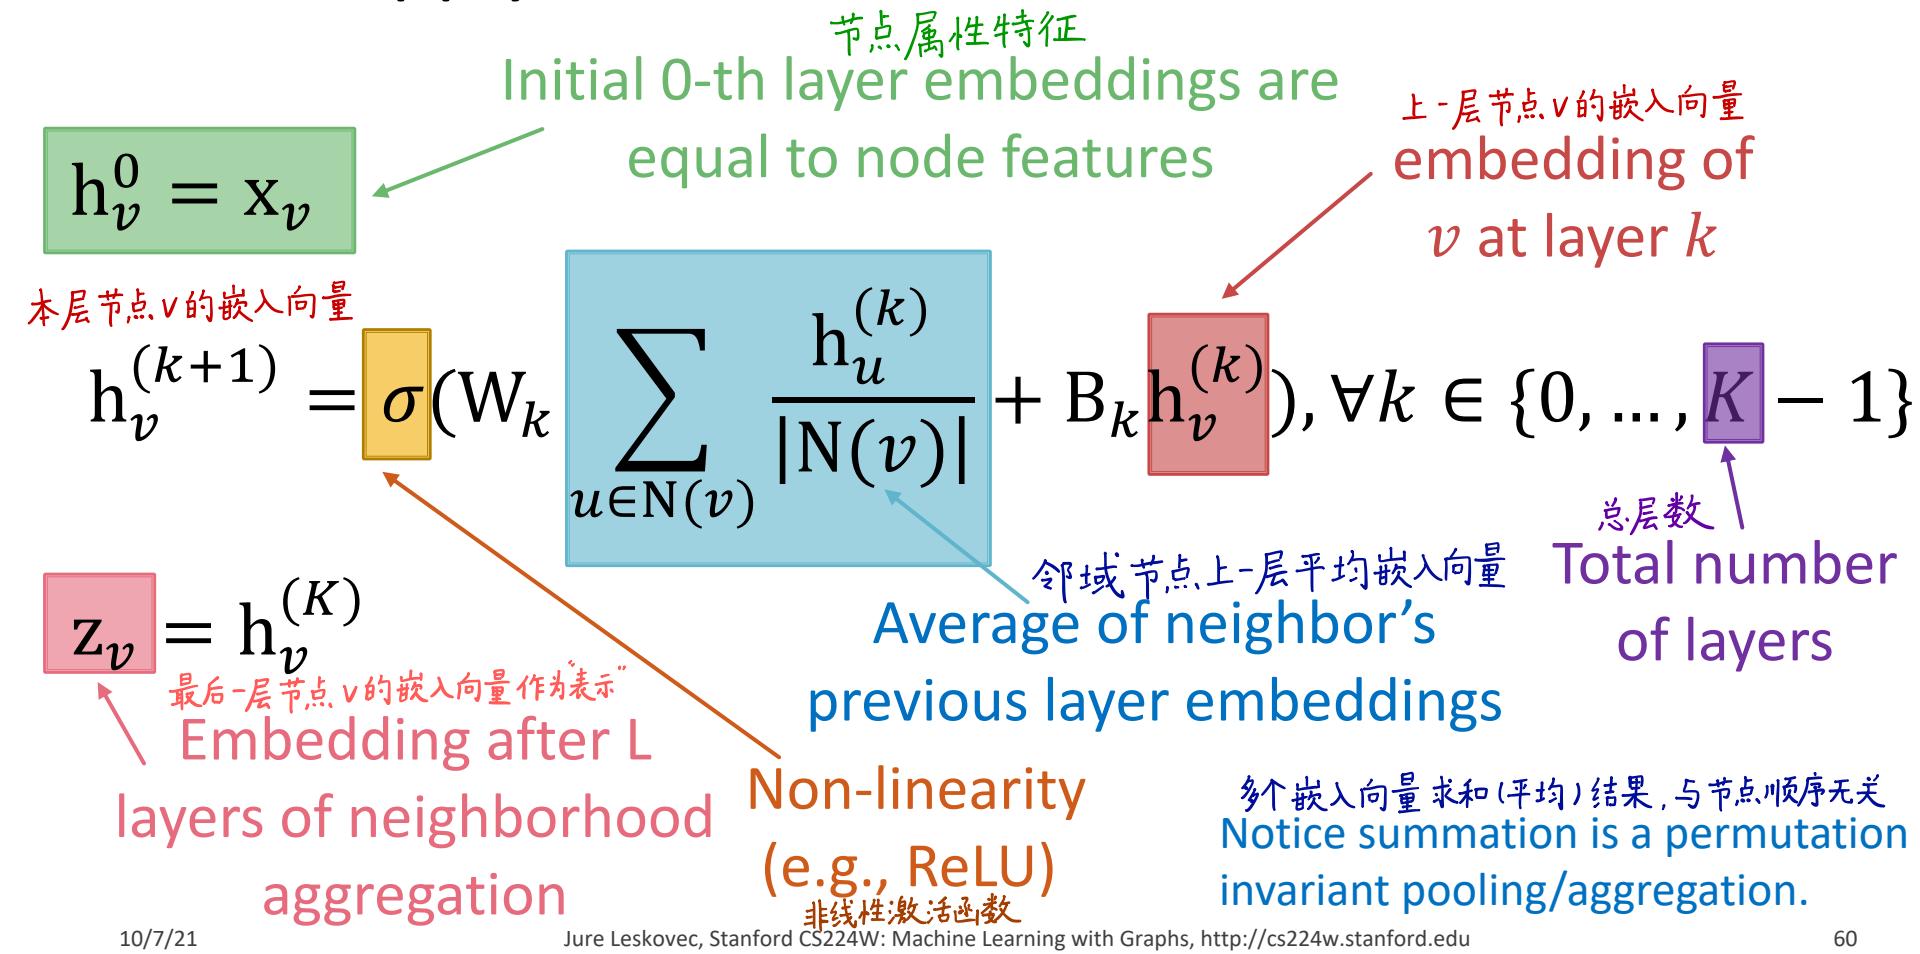
# Neighborhood Aggregation

- **Basic approach:** Average information from neighbors and apply a neural network



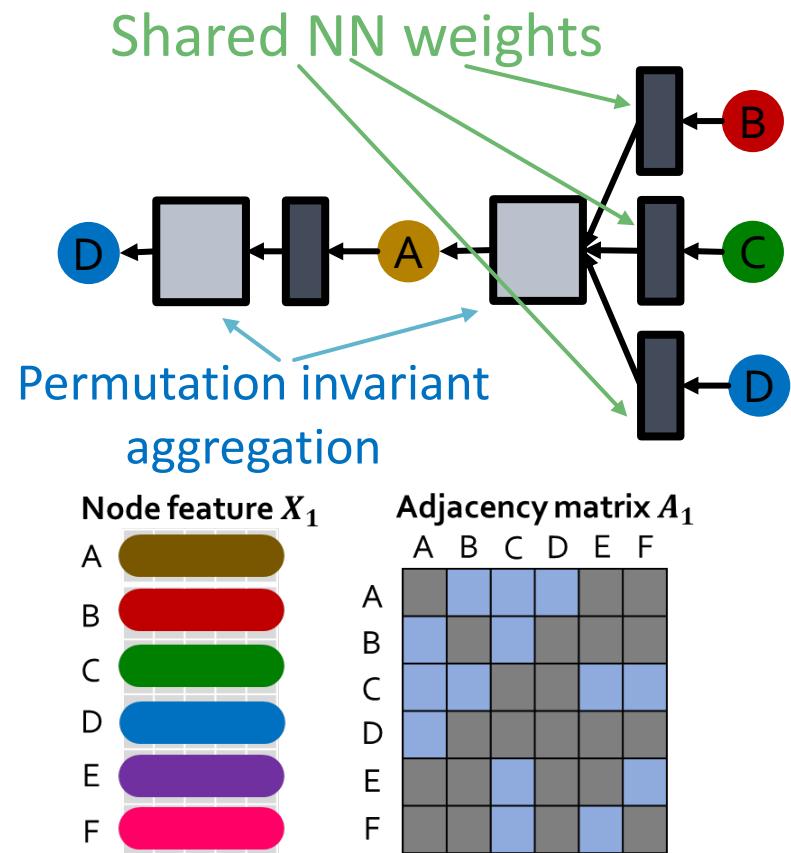
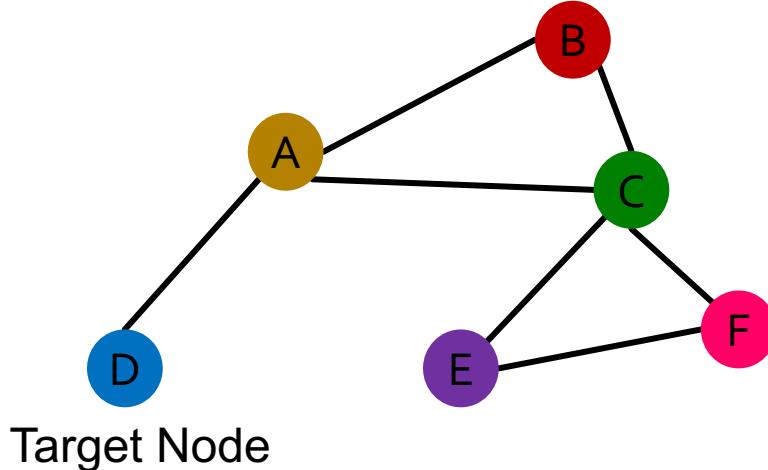
# The Math: Deep Encoder

- **Basic approach:** Average neighbor messages and apply a neural network



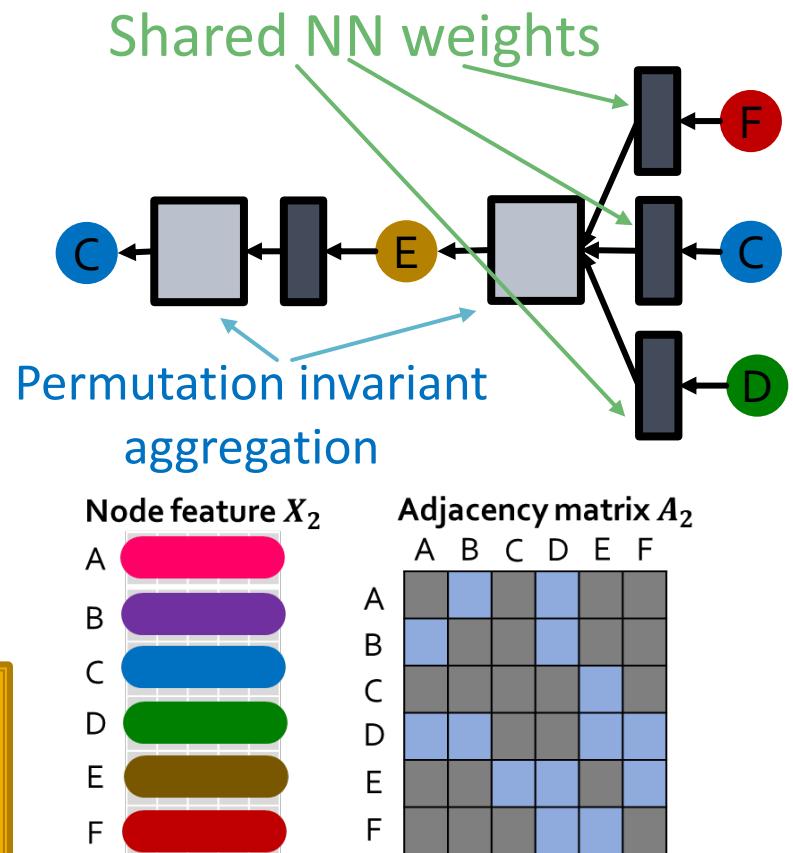
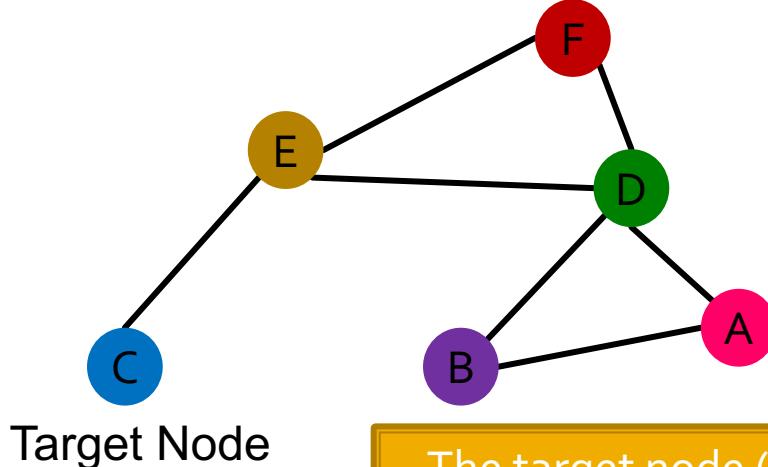
# Equivariant Property

Message passing and neighbor aggregation in graph convolution networks is permutation equivariant.



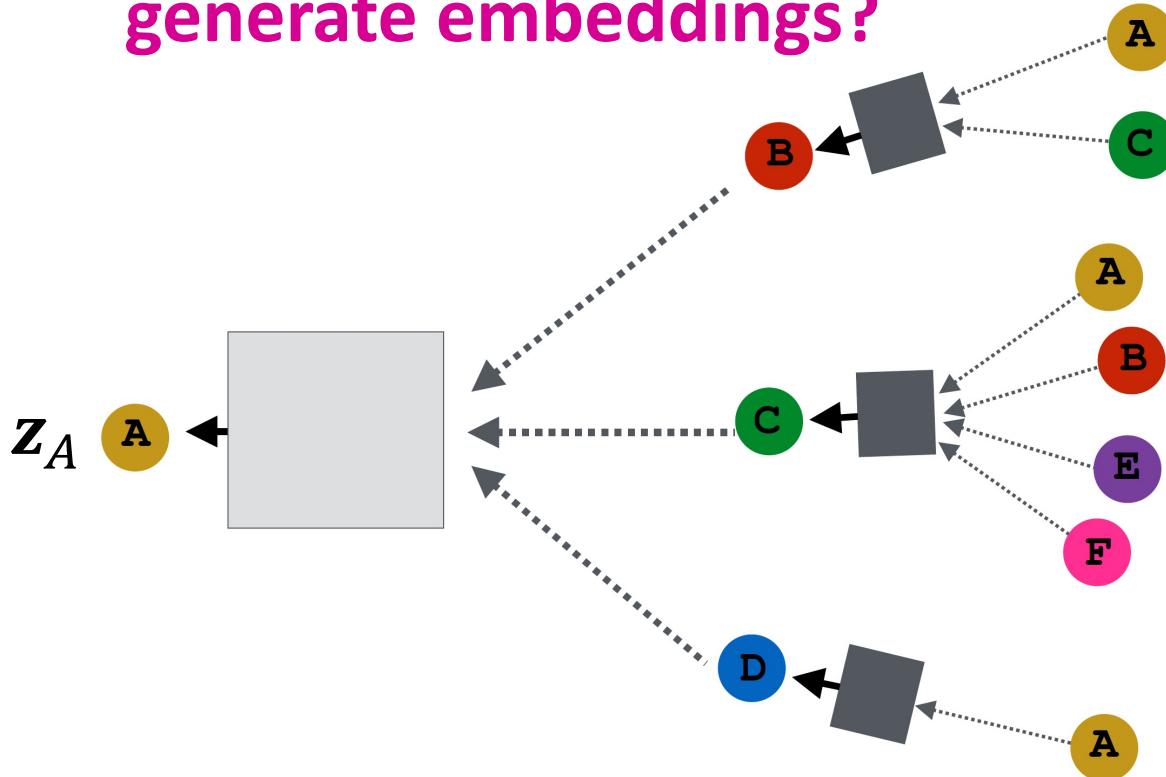
# Equivariant Property

Message passing and neighbor aggregation in graph convolution networks is permutation equivariant.



# Training the Model

How do we train the GCN to generate embeddings?



Need to define a loss function on the embeddings.

# Model Parameters

$$h_v^{(0)} = x_v$$

$$h_v^{(k+1)} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|} + B_k h_v^{(k)}), \forall k \in \{0..K-1\}$$

$$z_v = h_v^{(K)}$$

Trainable weight matrices

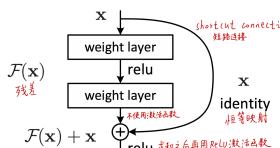
(i.e., what we learn)

第 $k$ 层图神经网络  
需训练学习得到的权重参数

$$\sum_{u \in N(v)} \frac{h_u^{(k)}}{|N(v)|}$$

非恒等映射  
不是残差连接

Final node embedding



We can feed these **embeddings** into any loss function and run SGD to **train the weight parameters**

$h_v^k$ : the hidden representation of node  $v$  at layer  $k$

- $W_k$ : weight matrix for neighborhood aggregation
- $B_k$ : weight matrix for transforming hidden vector of self

# Matrix Formulation (1)

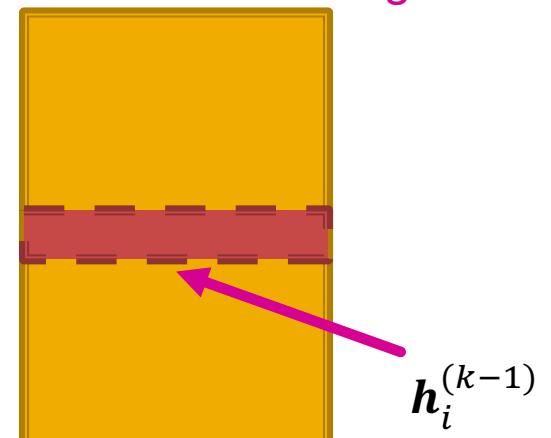
- Many aggregations can be performed efficiently by (sparse) matrix operations
- Let  $H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^T$
- Then:  $\sum_{u \in N_v} h_u^{(k)} = A_{v,:} H^{(k)}$
- Let  $D$  be diagonal matrix where  $D_{v,v} = \text{Deg}(v) = |N(v)|$ 
  - The inverse of  $D$ :  $D^{-1}$  is also diagonal:  
$$D_{v,v}^{-1} = 1/|N(v)|$$
- Therefore,

$$\sum_{u \in N(v)} \frac{h_u^{(k-1)}}{|N(v)|}$$



$$H^{(k+1)} = D^{-1} A H^{(k)}$$

Matrix of hidden embeddings  $H^{(k-1)}$

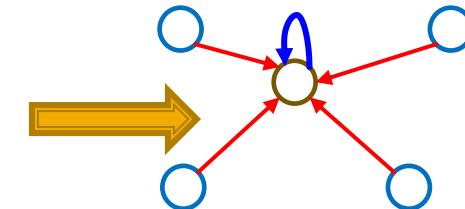


# Matrix Formulation (2)

- Re-writing update function in matrix form:

$$H^{(k+1)} = \sigma(\tilde{A}H^{(k)}W_k^T + H^{(k)}B_k^T)$$

where  $\tilde{A} = D^{-1}A$



$$H^{(k)} = [h_1^{(k)} \dots h_{|V|}^{(k)}]^T$$

- Red: neighborhood aggregation
- Blue: self transformation
- In practice, this implies that efficient sparse matrix multiplication can be used ( $\tilde{A}$  is sparse)
- **Note:** not all GNNs can be expressed in matrix form, when aggregation function is complex

# How to Train A GNN

- Node embedding  $\mathbf{z}_v$  is a function of input graph
- **Supervised setting:** we want to minimize the loss

$\mathcal{L}$  (see also Slide 15):  $f$ : 分类或回归预测头

$$\min_{\Theta} \mathcal{L}(y, f(\mathbf{z}_v))$$

- $y$ : node label 节点类别标注
- $\mathcal{L}$  could be L2 if  $y$  is real number, or cross entropy if  $y$  is categorical
- **Unsupervised setting:**
  - No node label available
  - Use the graph structure as the supervision! 自监督

# Unsupervised Training

- “Similar” nodes have similar embeddings

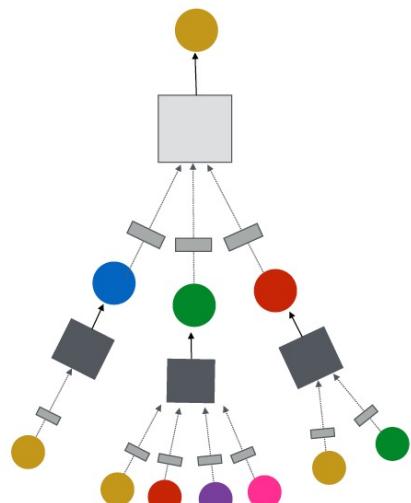
$$\mathcal{L} = \sum_{z_u, z_v} \text{CE}(y_{u,v}, \text{DEC}(z_u, z_v))$$

- Where  $y_{u,v} = 1$  when node  $u$  and  $v$  are **similar**
- **CE** is the cross entropy (Slide 16)
- **DEC** is the decoder such as inner product (Lecture 4)
- **Node similarity** can be anything from Lecture 3, e.g., a loss based on:
  - **Random walks** (node2vec, DeepWalk, struc2vec)
  - **Matrix factorization**
  - **Node proximity in the graph**

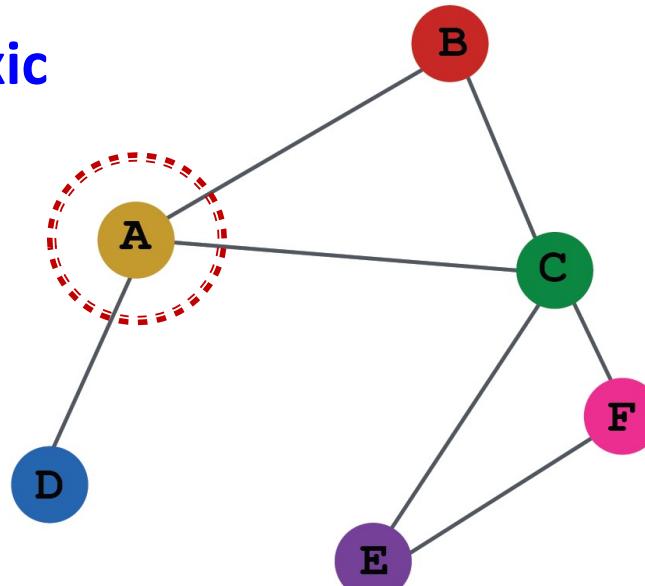
# Supervised Training

**Directly train** the model for a supervised task  
(e.g., node classification)

Safe or toxic  
drug?



Safe or toxic  
drug?



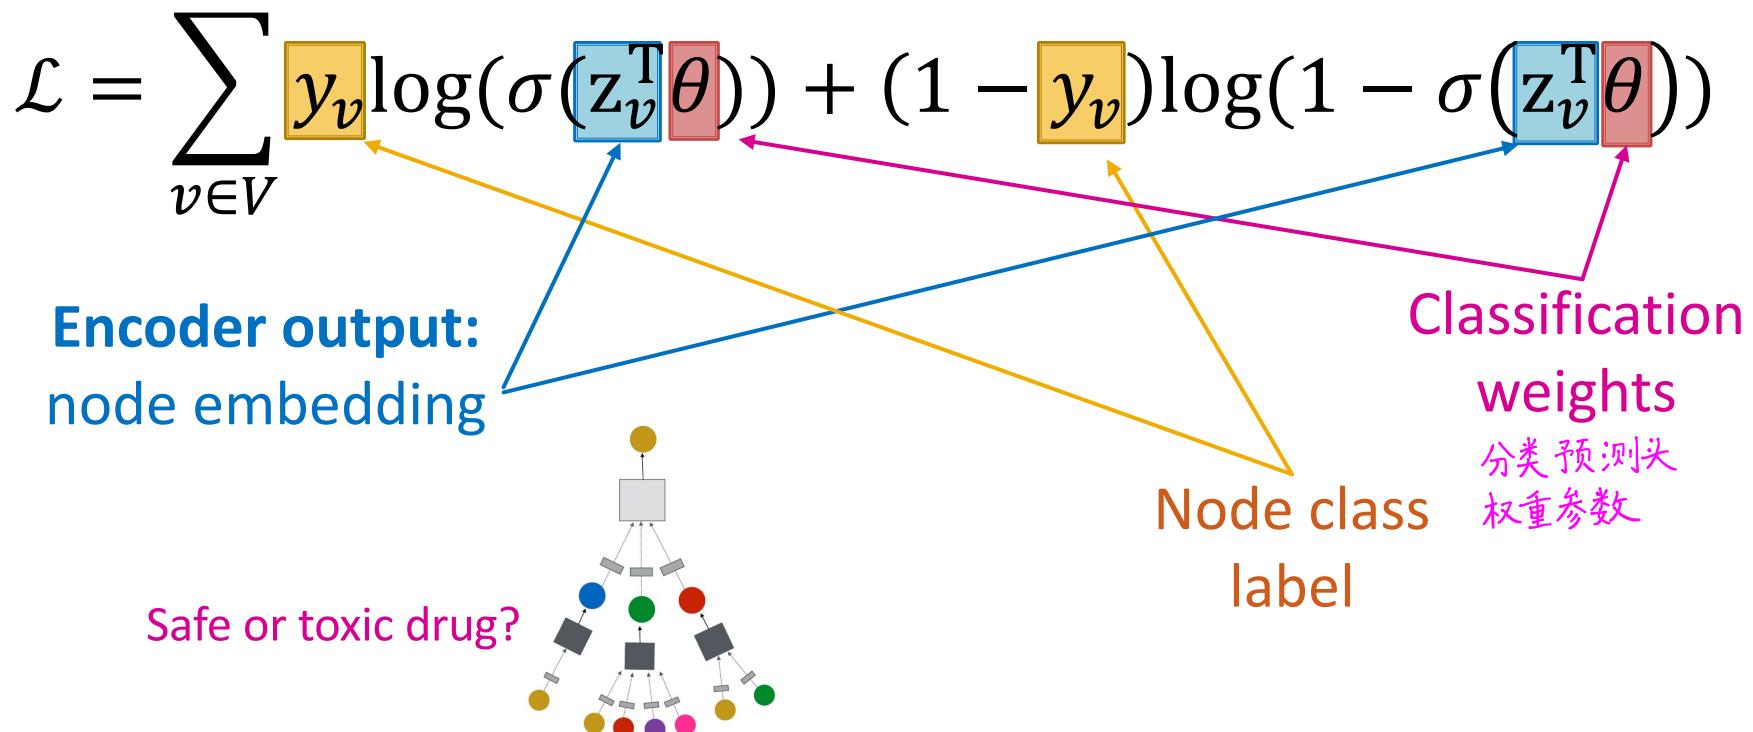
E.g., a drug-drug  
interaction network

药物-药物相互作用  
节点是药物

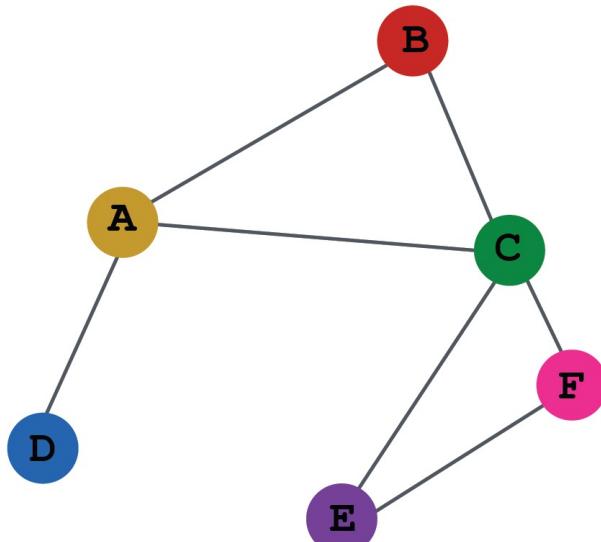
# Supervised Training

Directly train the model for a supervised task  
(e.g., node classification)

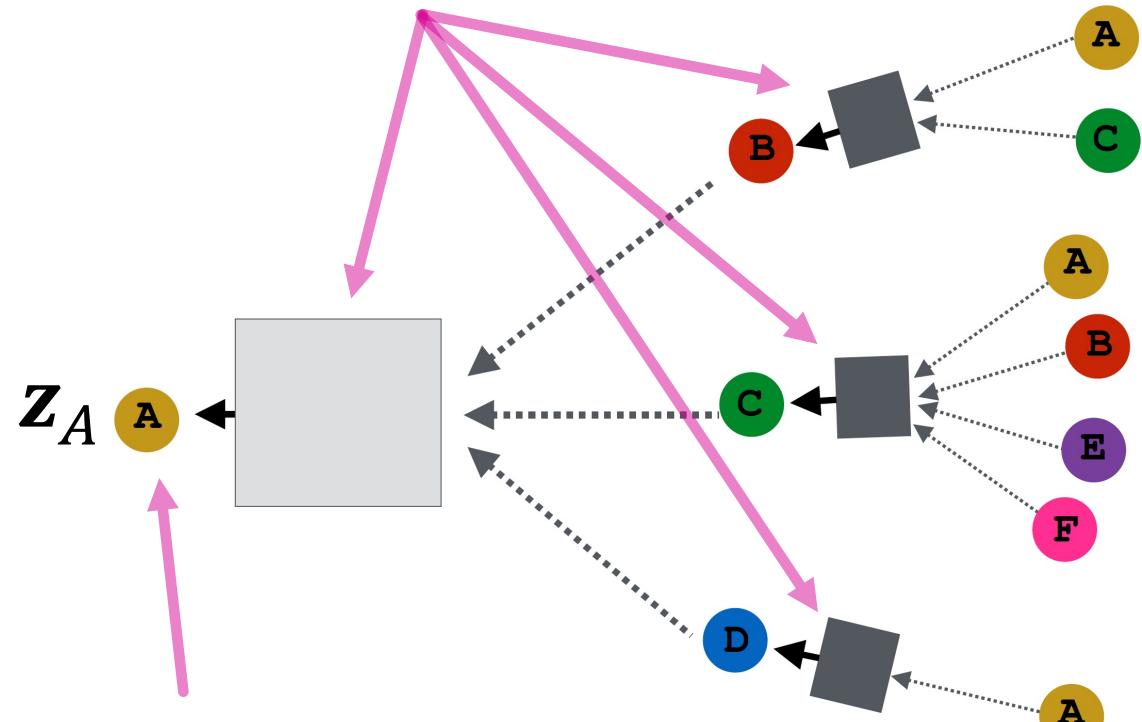
- Use cross entropy loss (Slide 16)



# Model Design: Overview

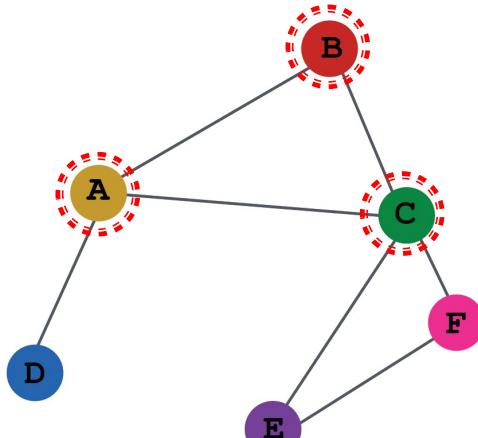


(1) Define a neighborhood aggregation function 聚合邻域信息的方式



(2) Define a loss function on the embeddings 损失函数

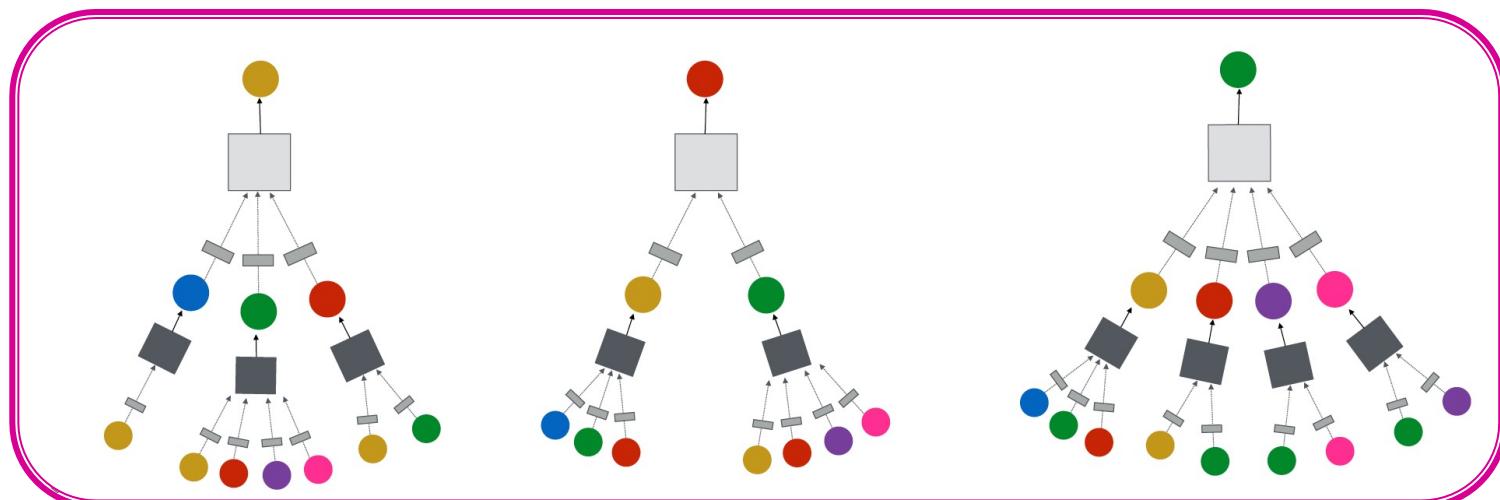
# Model Design: Overview



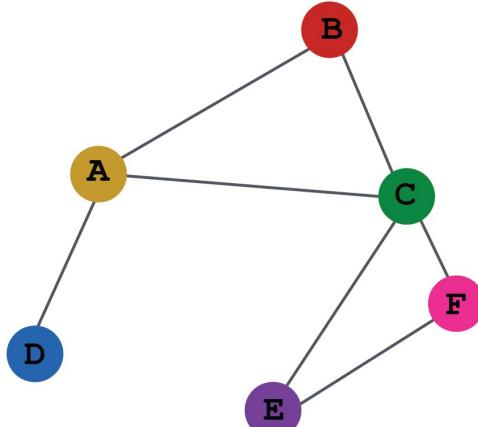
INPUT GRAPH

(3) Train on a set of nodes, i.e.,  
a batch of compute graphs

Mini batch 训练



# Model Design: Overview

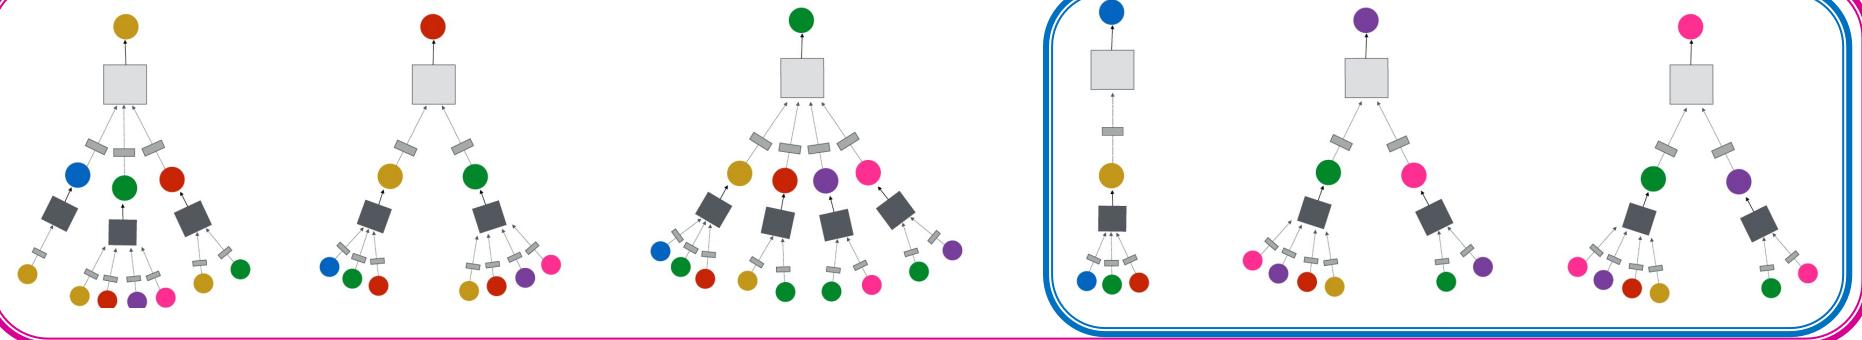


INPUT GRAPH

(4) Generate embeddings  
for nodes as needed

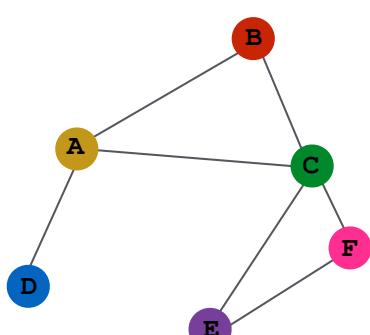
Even for nodes we never  
trained on!

泛化到新节点、新图  
Inductive Learning  
归纳式学习

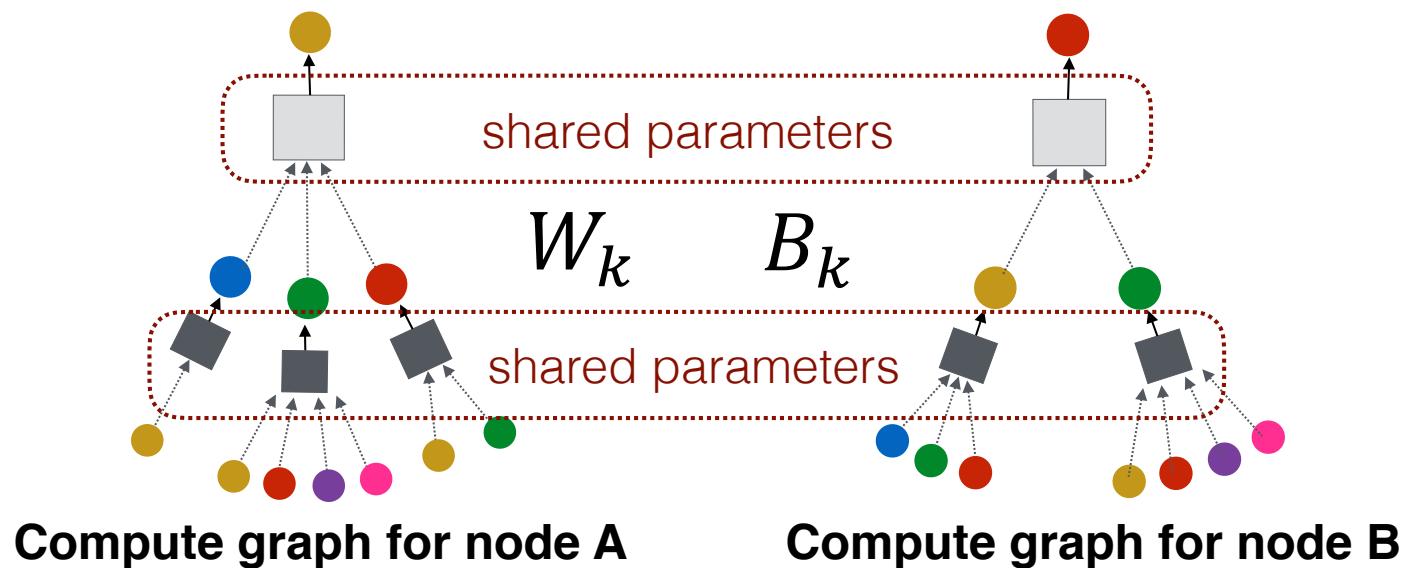


# Inductive Capability

- The same aggregation parameters are shared for all nodes:
  - The number of model parameters is sublinear in  $|V|$  and we can **generalize to unseen nodes!**

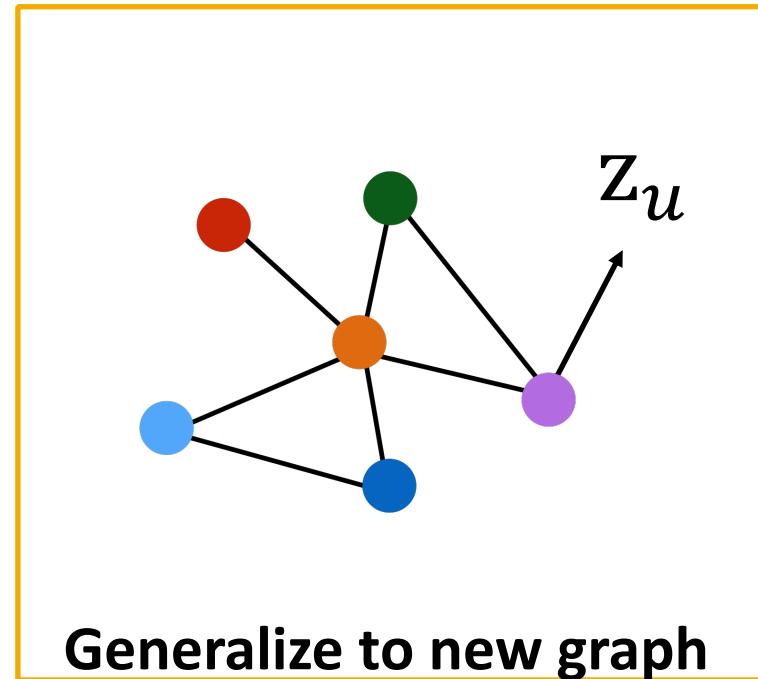
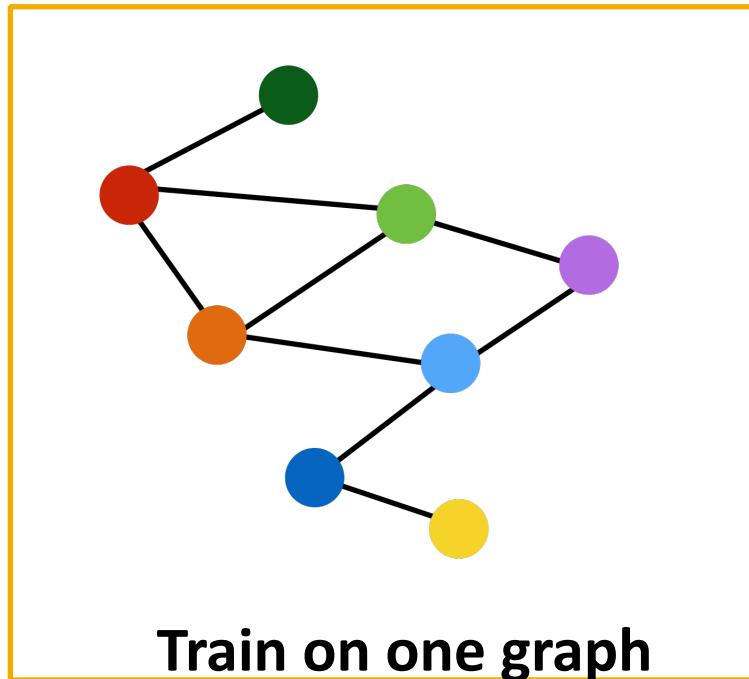


INPUT GRAPH



# Inductive Capability: New Graphs

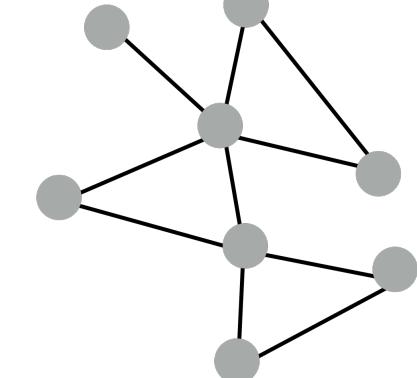
Inductive Learning 归纳式学习：泛化到新节点、甚至新图



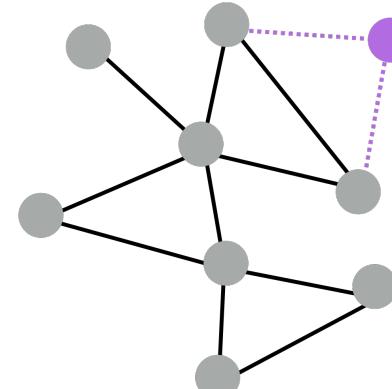
Inductive node embedding → Generalize to entirely unseen graphs

E.g., train on protein interaction graph from model organism A and generate embeddings on newly collected data about organism B  
组织

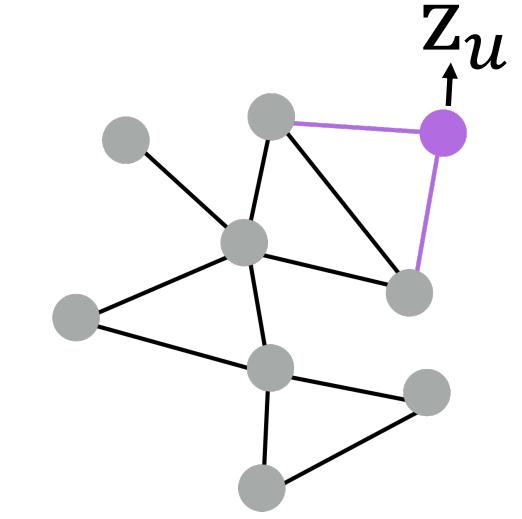
# Inductive Capability: New Nodes



Train with snapshot



New node arrives



Generate embedding  
for new node

- Many application settings constantly encounter previously unseen nodes:
  - E.g., Reddit, YouTube, Google Scholar
- Need to generate new embeddings “on the fly” 随用随取  
冷启动

# Outline of Today's Lecture

1. Basics of deep learning



2. Deep learning for graphs



3. Graph Convolutional Networks



4. GNNs subsume CNNs and  
Transformers

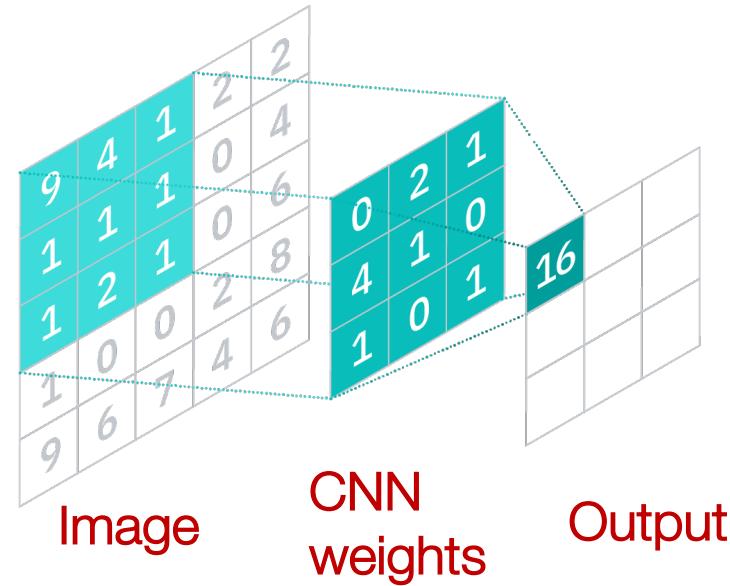
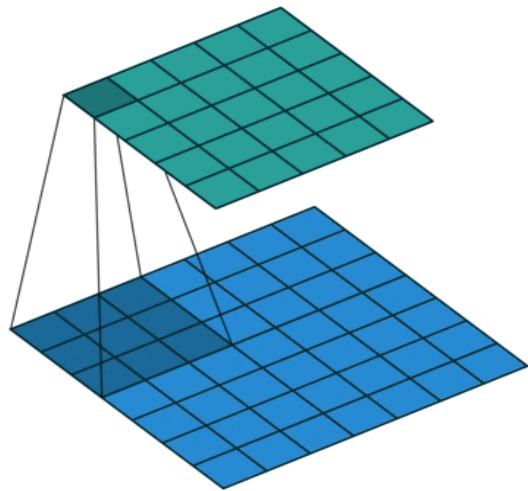


# Architecture Comparison

- How does GNNs compare to prominent architectures such as Convolutional Neural Nets, and Transformers?  
主要的

# Convolutional Neural Network

Convolutional neural network (CNN) layer with 3x3 filter:

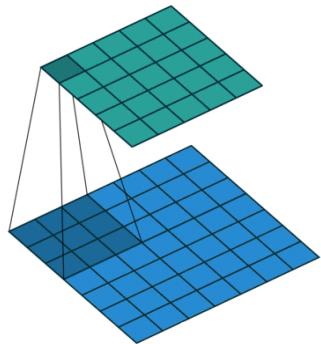


$$\text{CNN formulation: } h_v^{(l+1)} = \sigma(\sum_{u \in N(v) \cup \{v\}} W_l^u h_u^{(l)}), \quad \forall l \in \{0, \dots, L-1\}$$

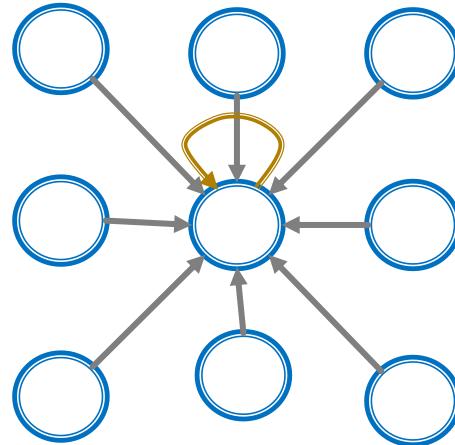
**$N(v)$  represents the 8 neighbor pixels of  $v$ .**

# GNN vs. CNN

Convolutional neural network (CNN) layer with 3x3 filter:



Image

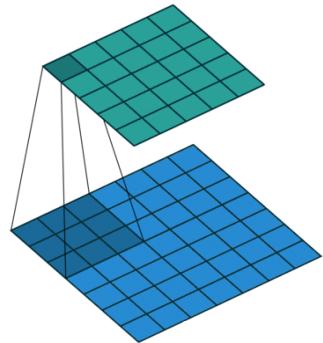


Graph

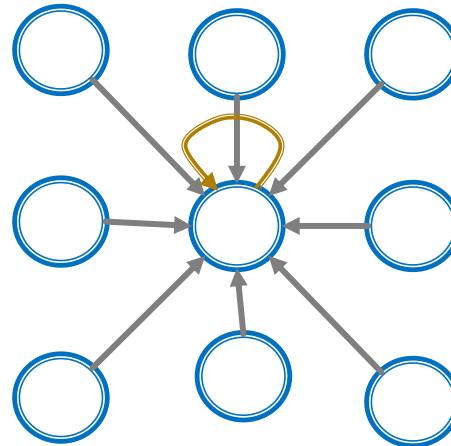
- GNN formulation (previous slide):  $h_v^{(l+1)} = \sigma(\mathbf{W}_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)})$ ,  $\forall l \in \{0, \dots, L-1\}$
- CNN formulation:  
if we rewrite:  
$$h_v^{(l+1)} = \sigma(\sum_{u \in N(v) \cup \{v\}} W_l^u h_u^{(l)})$$
$$h_v^{(l+1)} = \sigma(\sum_{u \in N(v)} \mathbf{W}_l^u h_u^{(l)} + B_l h_v^{(l)})$$
,  $\forall l \in \{0, \dots, L-1\}$

# GNN vs. CNN

Convolutional neural network (CNN) layer with 3x3 filter:



Image



Graph

$$\text{GNN formulation: } h_v^{(l+1)} = \sigma(\mathbf{W}_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$$

$$\text{CNN formulation: } h_v^{(l+1)} = \sigma(\sum_{u \in N(v)} \mathbf{W}_l^u h_u^{(l)} + B_l h_v^{(l)}), \forall l \in \{0, \dots, L-1\}$$

**Key difference:** We can learn different  $W_l^u$  for different “neighbor”  $u$  for pixel  $v$  on the image. The reason is we can pick an order for the 9 neighbors using **relative position** to the center pixel:  $\{(-1, -1), (-1, 0), (-1, 1), \dots, (1, 1)\}$

# GNN vs. CNN

Convolutional neural network (CNN) layer with 3x3 filter:



固定的邻域和固定的顺序

- CNN can be seen as a special GNN with **fixed neighbor size and ordering**:  
    • The size of the filter is pre-defined for a CNN.  
    • The advantage of GNN is it processes arbitrary graphs with different degrees for each node.

卷积核

任意的

**Key difference:** We can learn different  $W_l^u$  for different “neighbor”  $u$  for pixel  $v$  on the image. The reason is we can pick an order for the 9 neighbors using **relative position** to the center pixel:  $\{(-1, -1), (-1, 0), (-1, 1), \dots, (1, 1)\}$

# GNN vs. CNN

Convolutional neural network (CNN) layer with 3x3 filter:

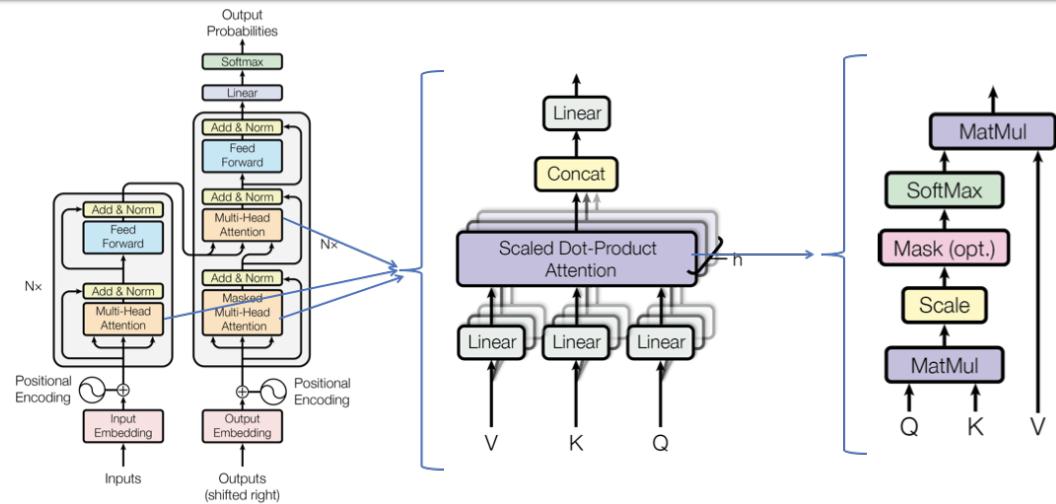


- CNN can be seen as a special GNN with fixed neighbor size and ordering.
- CNN is **not** permutation equivariant.
  - Switching the order of pixels will leads to different outputs.

**Key difference:** We can learn different  $W_l^u$  for different “neighbor”  $u$  for pixel  $v$  on the image. The reason is we can pick an order for the 9 neighbors using **relative position** to the center pixel:  $\{(-1, -1), (-1, 0), (-1, 1), \dots, (1, 1)\}$

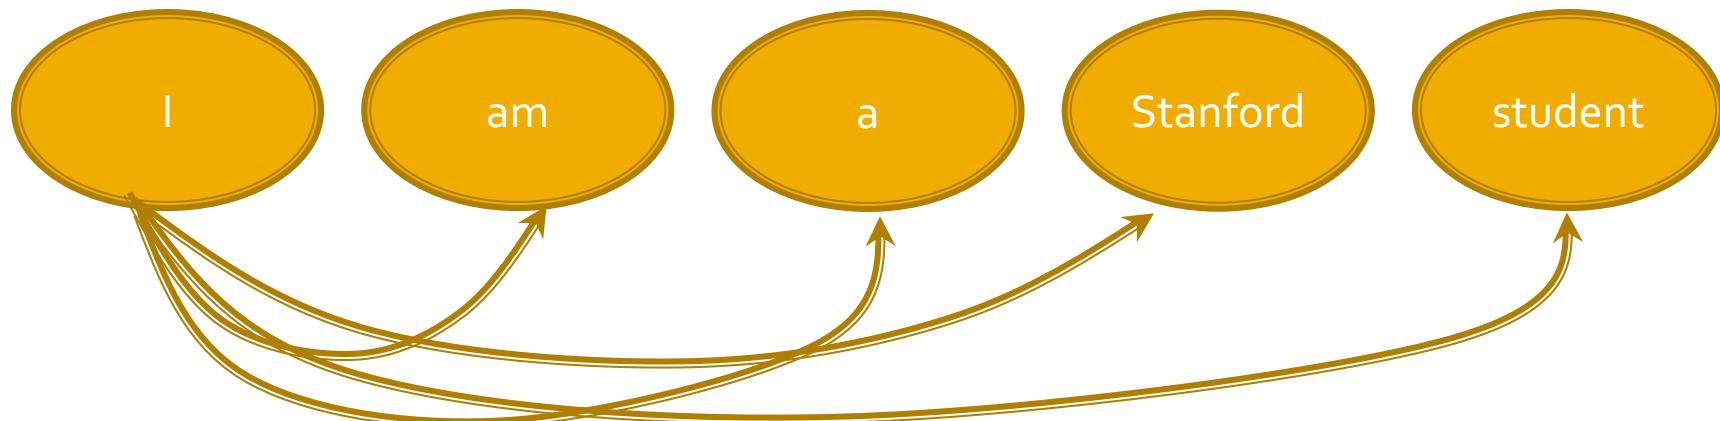
# Transformer

Transformer is one of the most popular architectures that achieves great performance in many sequence modeling tasks.



**Key component: self-attention** 自注意力

- Every token/word attends to all the other tokens/words via matrix calculation.

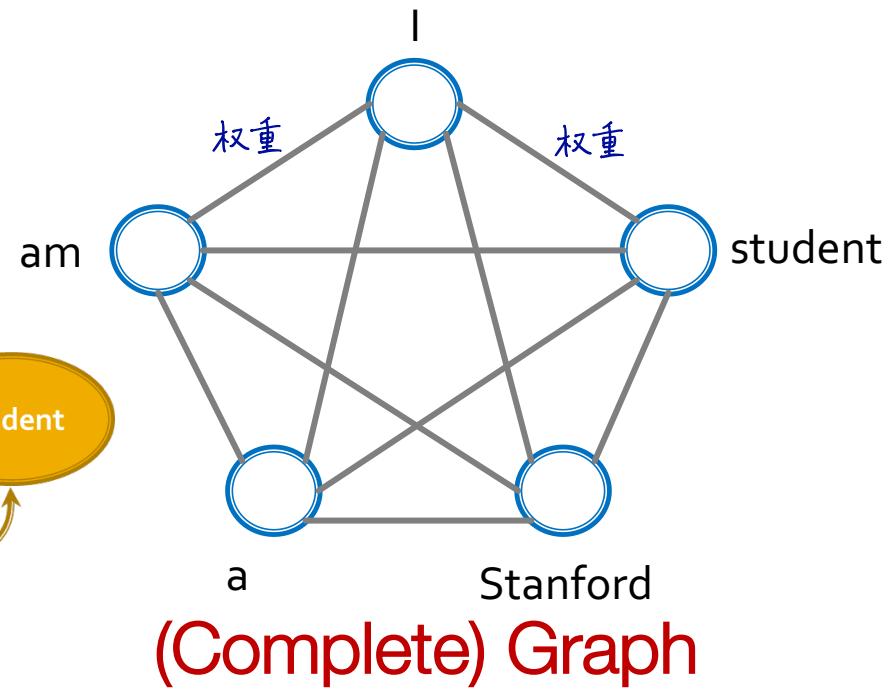
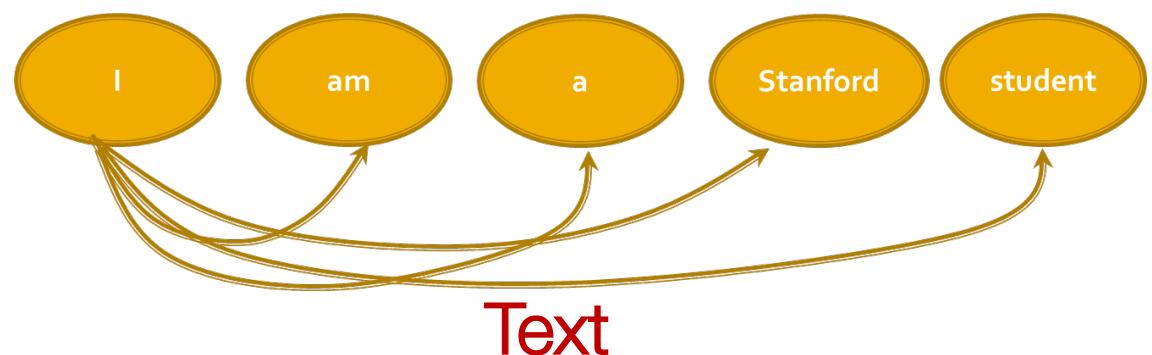


# GNN vs. Transformer

Transformer layer can be seen as a special GNN that runs on a fully-connected “word” graph!

Graph Attention Network (GAT)

Since each word attends to **all the other words**, **the computation graph** of a transformer layer is identical to that of a GNN on the **fully-connected “word” graph**.



# Summary

- In this lecture, we introduced
  - Basics of neural networks
    - Loss, Optimization, Gradient, SGD, non-linearity, MLP
  - Idea for Deep Learning for Graphs
    - Multiple layers of embedding transformation
    - At every layer, use the embedding at previous layer as the input
    - Aggregation of neighbors and self-embeddings
  - Graph Convolutional Network 图卷积神经网络
    - Mean aggregation; can be expressed in matrix form
  - GNN is a general architecture
    - CNN and Transformer can be viewed as a special GNN

计算图

